



LINUX
SECURITY
SUMMIT
EUROPE

Compartmentalizing Vulnerable Kernel Components Without Stopping the Machine

Presenter: Qinrun Dai

Contributors: Zicheng Wang, Tiejin Chen, Yueqi Chen, and Hua Wei

About us




Qinrun Dai

PhD Student
University of Colorado, Boulder



Zicheng Wang

PhD
Nanjing University



Tiejin Chen

PhD Student
Arizona State University



Yequi Chen

Assistant Professor
University of Colorado, Boulder



Hua Wei

Assistant Professor
Arizona State University

Agenda

- **Motivation**
- Challenges
- Overview & Technical Details
- Demo (CVE-2022-0995)
- Evaluation

Tool is available at:

<https://github.com/a8stract-lab/o2c>

Paper is available at:

<https://arxiv.org/abs/2401.05641>

Risky Time Window in Kernel Development



Risky Time Window in Kernel Development



Linux Kernel Development Timeline



Risky Time Window in Kernel Development



Risky Time Window in Kernel Development



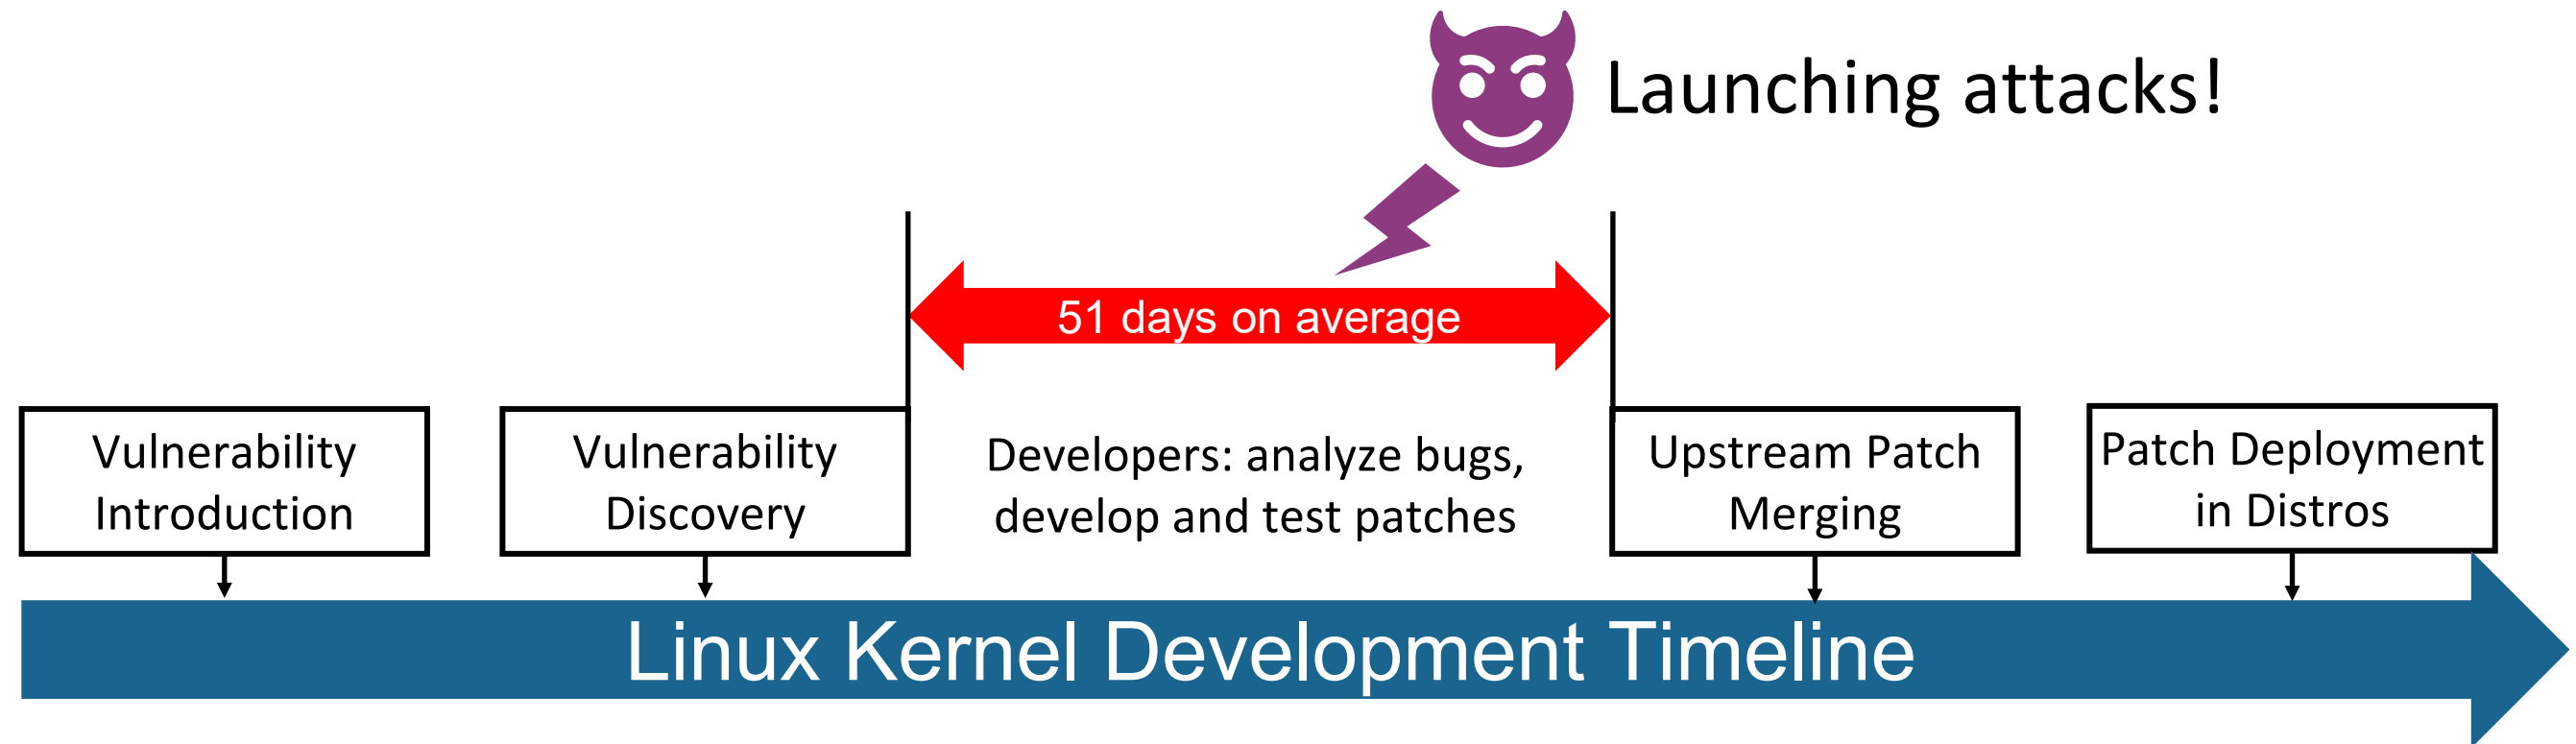
Risky Time Window in Kernel Development



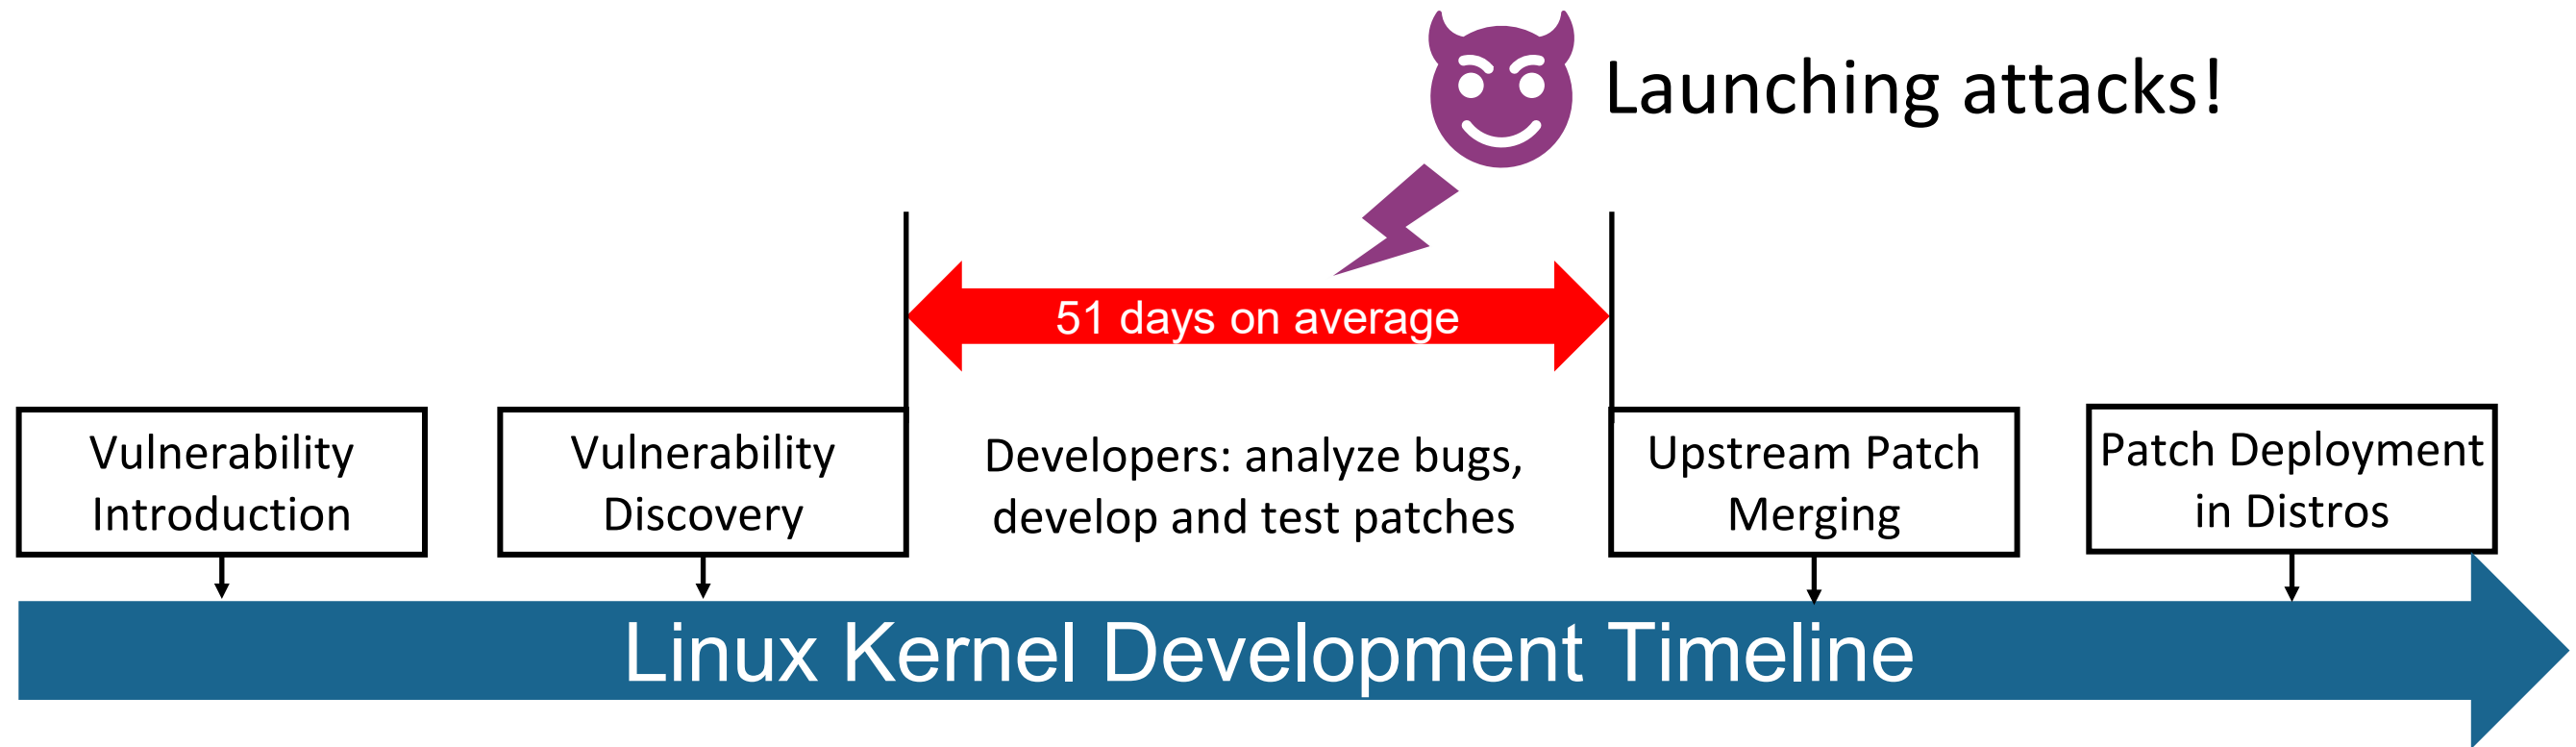
Risky Time Window in Kernel Development



Risky Time Window in Kernel Development




Risky Time Window in Kernel Development




How to remediate newly discovered vulnerabilities before official patches are available?

Disruptive Solution is Unacceptable



Disruptive Solution is Unacceptable




Disruptive Solution is Unacceptable

Takeaway:

A disruptive solution that requires rebooting and disrupting running service is unacceptable.

Otherwise over 3 times of rebooting is needed to have a full coverage.




Disruptive Solution is Unacceptable

Takeaway:

A disruptive solution that requires rebooting and disrupting running service is unacceptable.

Otherwise over 3 times of rebooting is needed to have a full coverage.

**An On-the-Fly
solution
is desired**

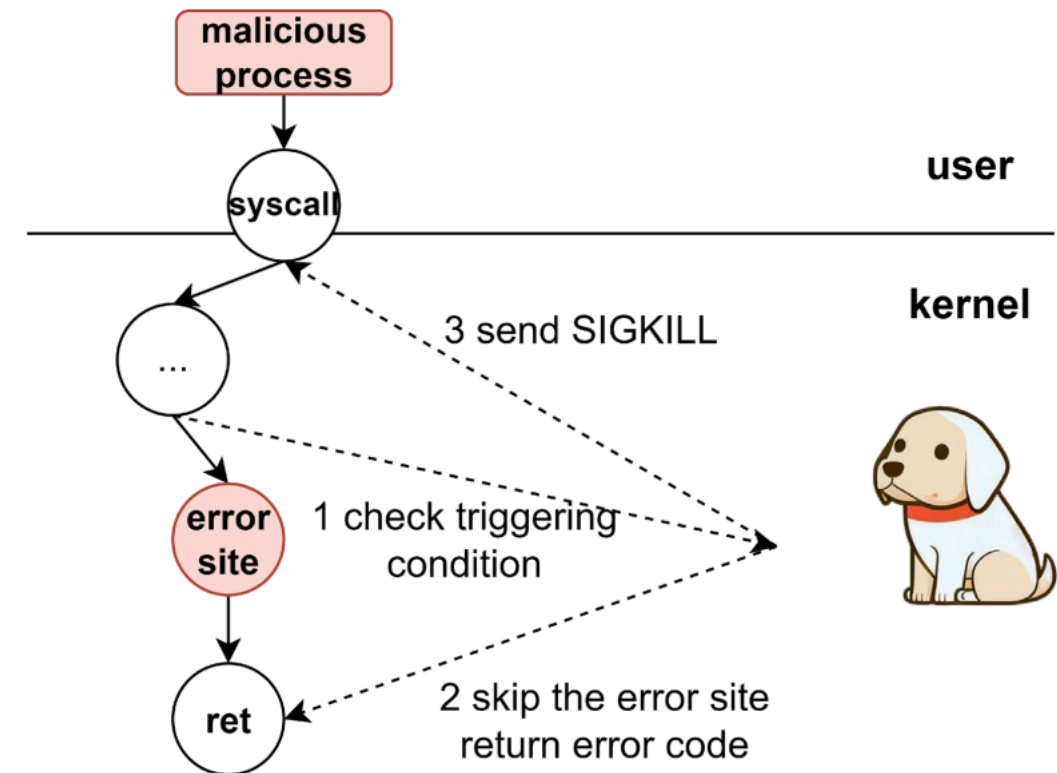


Existing Solutions and Limitations

[1] PET: Prevent Discovered Errors from Being Triggered in the Linux Kernel, USENIX Security'23

Existing Solutions and Limitations

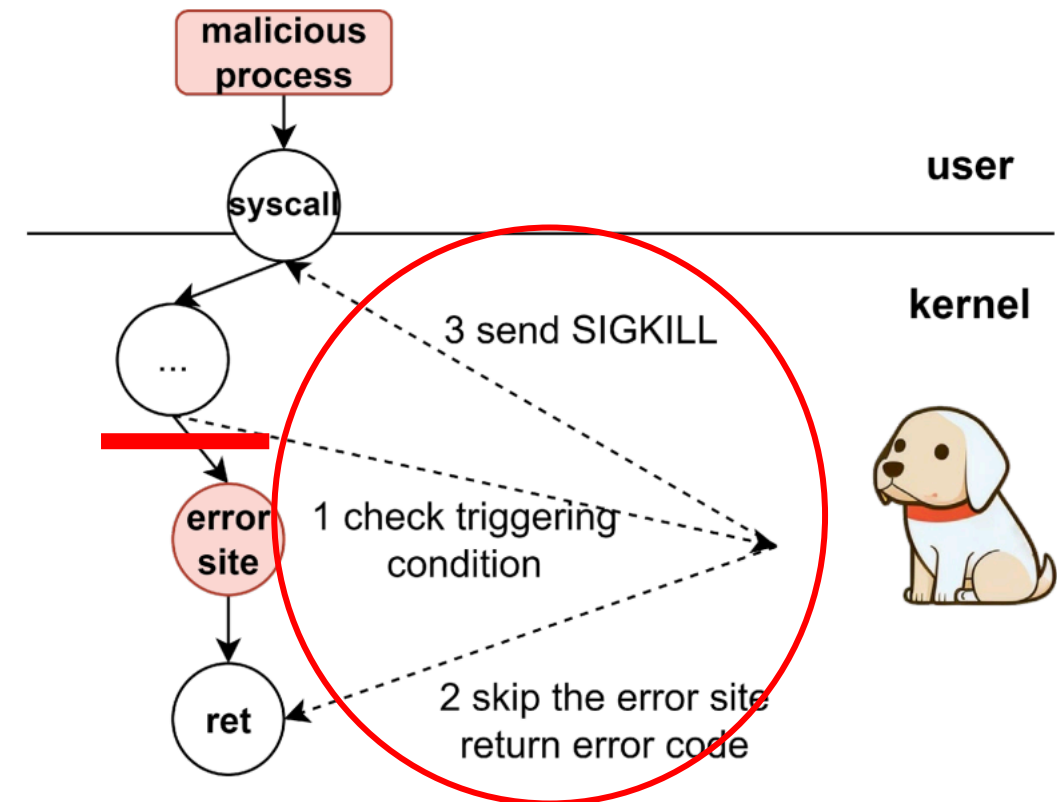
- PET^[1]



[1] PET: Prevent Discovered Errors from Being Triggered in the Linux Kernel, USENIX Security'23

Existing Solutions and Limitations

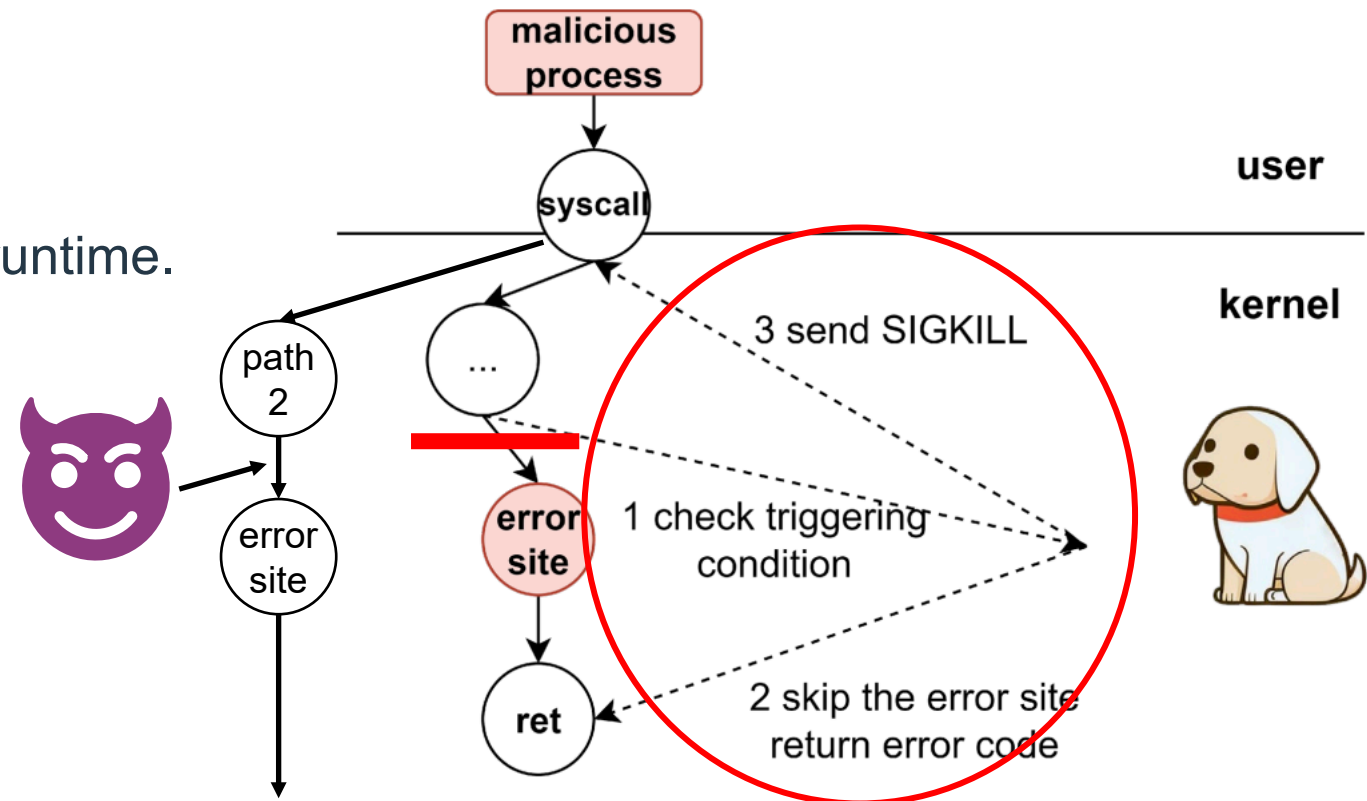
- PET^[1]
- Core idea:
 - Construct triggering conditions.
 - Determine if triggering condition is met at runtime.
 - Prevent triggering if yes.



[1] PET: Prevent Discovered Errors from Being Triggered in the Linux Kernel, USENIX Security'23

Existing Solutions and Limitations

- PET^[1]
- Core idea:
 - Construct triggering conditions.
 - Determine if triggering condition is met at runtime.
 - Prevent triggering if yes.
- Limitation:
 - Can be bypassed if exploits target another triggering site along a different path.



[1] PET: Prevent Discovered Errors from Being Triggered in the Linux Kernel, USENIX Security'23

Existing Solutions and Limitations (cont.)

[2] SeaK: Rethinking the Design of a Secure Allocator for OS Kernel, USENIX Security'24

Existing Solutions and Limitations (cont.)

- SeaK^[2]

[2] SeaK: Rethinking the Design of a Secure Allocator for OS Kernel, USENIX Security'24

Existing Solutions and Limitations (cont.)

- SeaK^[2]

vuln obj vic. obj

Typical memory layout of heap
out-of-bound exploitation

[2] SeaK: Rethinking the Design of a Secure Allocator for OS Kernel, USENIX Security'24

Existing Solutions and Limitations (cont.)

- SeaK^[2]
- Core idea:
 - Isolates vulnerable objects, victim objects, and spray objects in different regions.

vuln obj vic. obj

Typical memory layout of heap
out-of-bound exploitation

guard vic. obj guard vic. obj

guard vuln. obj guard vuln. obj

guard sprayed vic2 guard sprayed vic2

Memory layout after isolation

[2] SeaK: Rethinking the Design of a Secure Allocator for OS Kernel, USENIX Security'24

Existing Solutions and Limitations (cont.)

- SeaK^[2]
- Core idea:
 - Isolates vulnerable objects, victim objects, and spray objects in different regions.
- Limitation:
 - While more general than PET, SeaK^[2] can be bypassed if attackers exploit legacy objects.



Typical memory layout of heap out-of-bound exploitation



Memory layout after isolation


[2] SeaK: Rethinking the Design of a Secure Allocator for OS Kernel, USENIX Security'24

Legacy Objects Problem in Detail



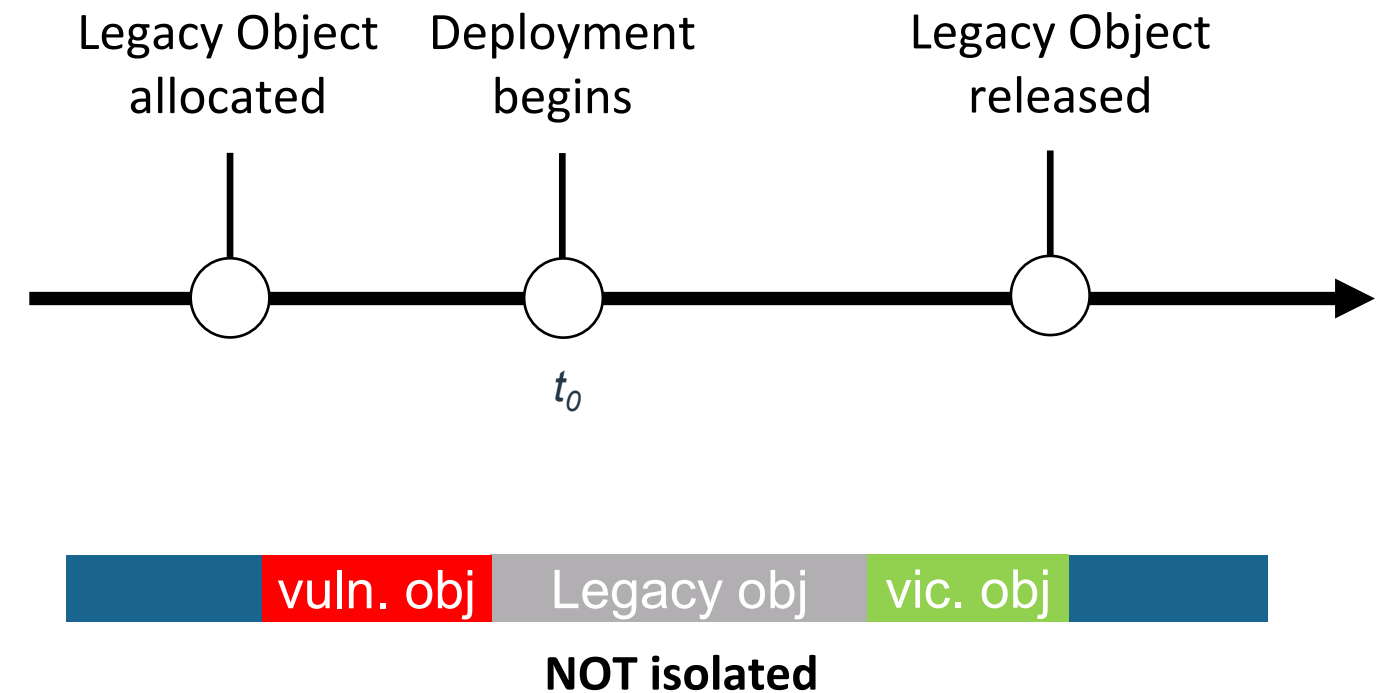
Legacy Objects Problem in Detail

- Definition: objects allocated before protection is deployed (t_0) and released after t_0 .




Legacy Objects Problem in Detail

- Definition: objects allocated before protection is deployed (t_0) and released after t_0 .
- What if a vulnerable / victim object is legacy?
 - Not isolated and mixed up with other objects.



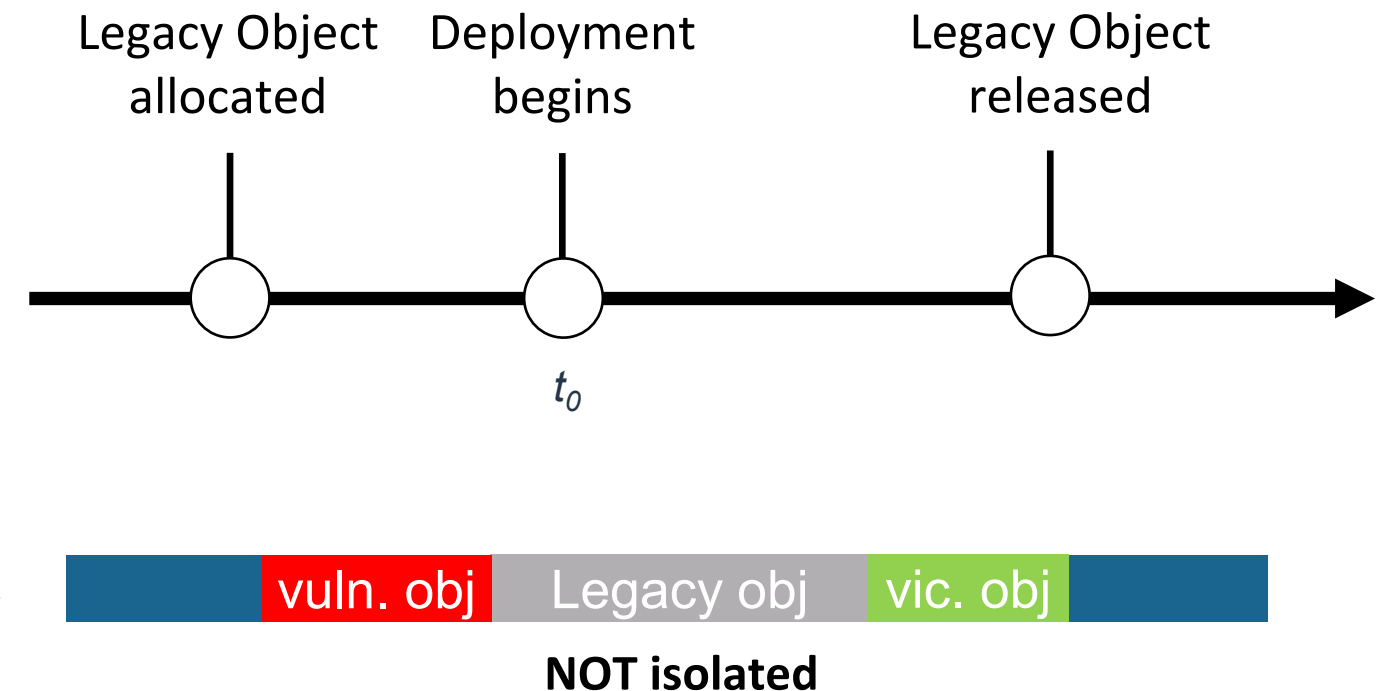
Legacy Objects Problem in Detail

- Definition: objects allocated before protection is deployed (t_0) and released after t_0 .
- What if a vulnerable / victim object is legacy?
 - Not isolated and mixed up with other objects.
- Many abuse chances:
 - Lifetime of legacy objects is long: 10,862 objects last more than 10s.
 - Many chances to manipulate legacy objects: average 22.87 modifications during the object's lifetime.



Legacy Objects Problem in Detail

- Definition: objects allocated before protection is deployed (t_0) and released after t_0 .
- What if a vulnerable / victim object is legacy?
 - Not isolated and mixed up with other objects.
- Many abuse chances:
 - Lifetime of legacy objects is long: 10,862 objects last more than 10s.
 - Many chances to manipulate legacy objects: average 22.87 modifications during the object's lifetime.



Auditing legacy objects access is the focus of this briefing

Agenda

- Motivation
- **Challenges**
- Overview & Technical Details
- Demo (CVE-2022-0995)
- Evaluation





Legacy Object Auditing - Challenge 1





Legacy Object Auditing - Challenge 1

- Fact: Legacy objects are allocated before protection is enabled. We cannot record KASAN-like metadata for legacy objects.



Legacy Object Auditing - Challenge 1

- Fact: Legacy objects are allocated before protection is enabled. We cannot record KASAN-like metadata for legacy objects.
- Consequence: When a legacy object is accessed, start address, end address, and type are untracked. Thus we don't know whether access is illegal or not.




Solution to Challenge 1



Solution to Challenge 1

C: `__set_bit(q->type, watch_filter->type_filter);`

Asm: `BTS [R15], RAX`




Solution to Challenge 1

- We use Machine Learning to infer the type of an accessed object, compared with access pointer type.

C: `__set_bit(q->type, watch_filter->type_filter);`


Asm: `BTS [R15], RAX`



Solution to Challenge 1

- We use Machine Learning to infer the type of an accessed object, compared with access pointer type.

```
C:      __set_bit(q->type, watch_filter->type_filter);
Asm:    BTS [R15], RAX
```



Human: What does these unorganized data mean?



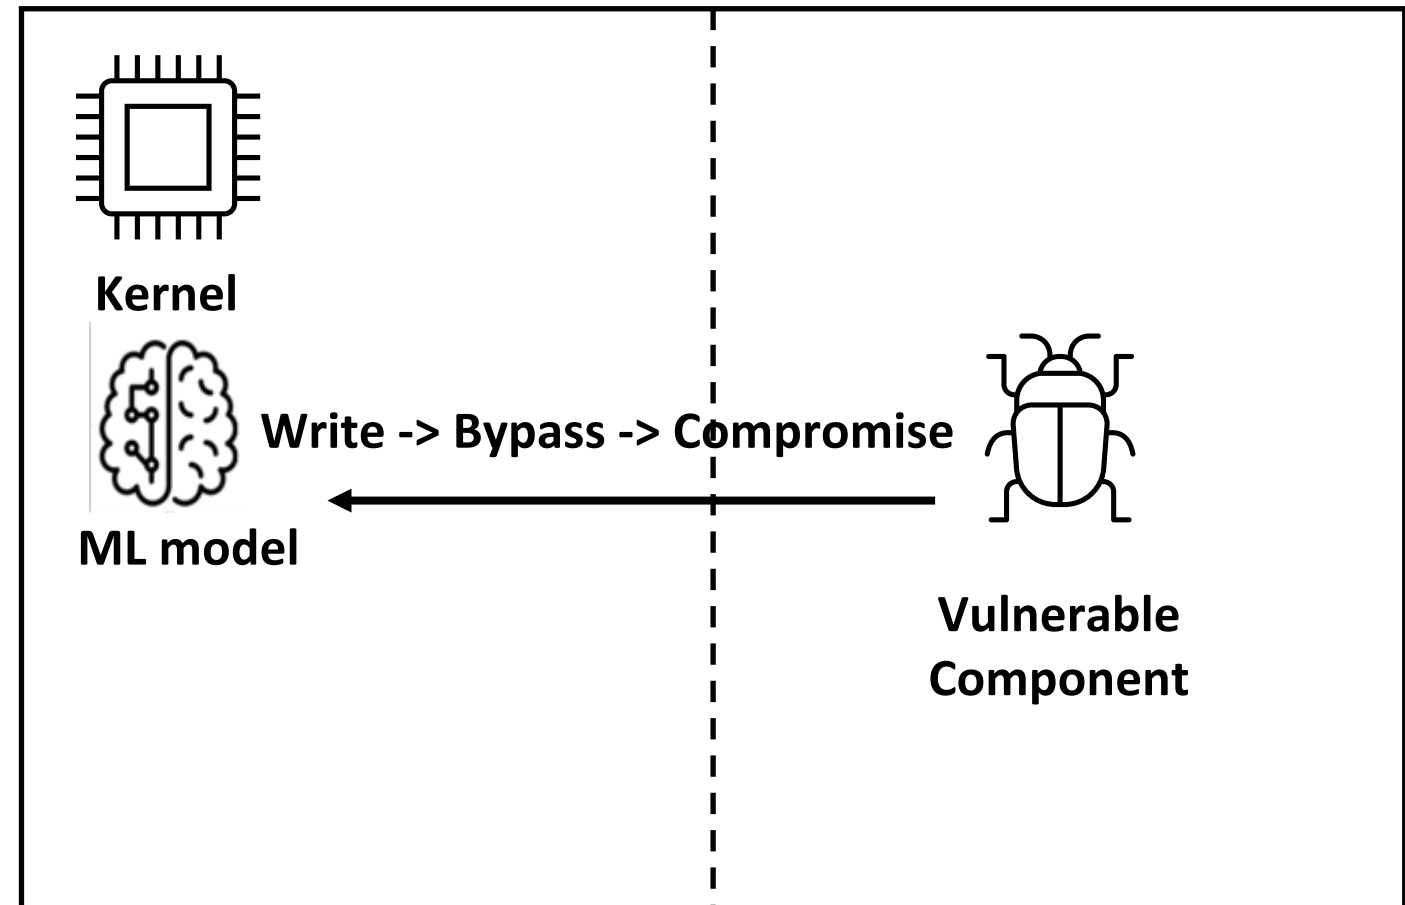
Trained AI: According to byte1, byte2, ..., byteN, the object's type is inferred as msg_msg, indicating error because expected type should be watch_filter.

0xffff88810738e5c0	41 62 73 74 72 61 63 74
0xffff88810738e5c8	A0 79 04 02 81 88 FF FF
0xffff88810738e5d0	00 AC 04 02 81 88 FF FF

0xffff88810738e5c0	41 62 73 74 72 61 63 74
0xffff88810738e5c8	A0 79 04 02 81 88 FF FF
0xffff88810738e5d0	00 AC 04 02 81 88 FF FF

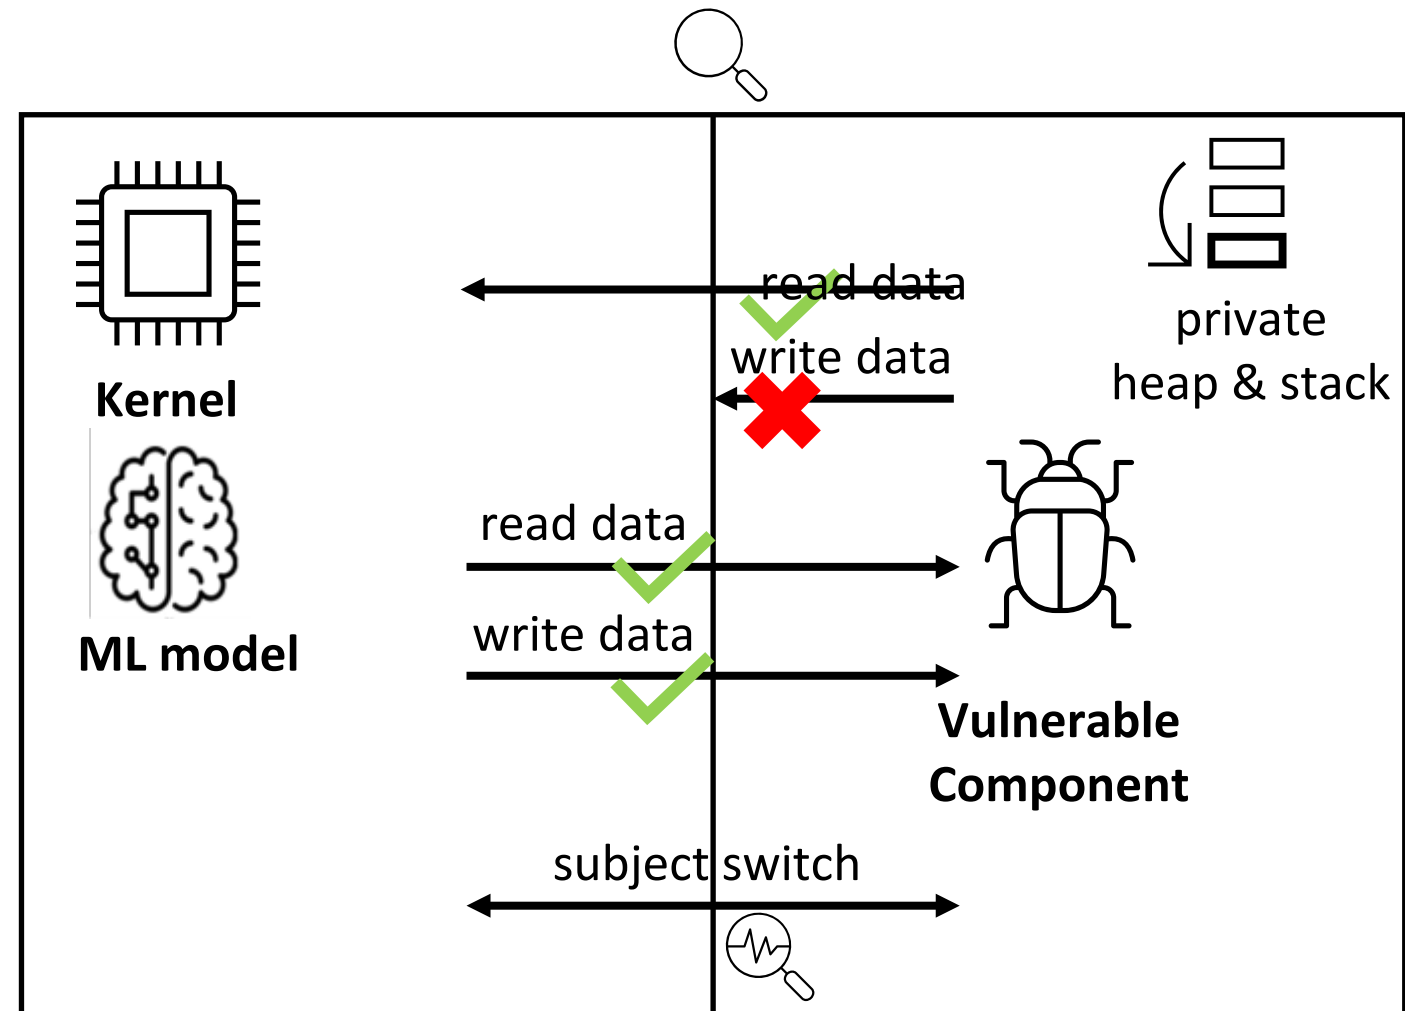
Legacy Object Auditing - Challenge 2

- Auditing integrity
- How to ensure the following integrity of auditing will not be compromised?
 - ML model integrity
 - Data-Flow integrity
 - Control-Flow integrity



Solution to Challenge 2

- Kernel Code instrumentation
 - Audit each read / write
 - Audit subject switch
- Private heap & stack
 - Vulnerable Component only use its own private data structures.

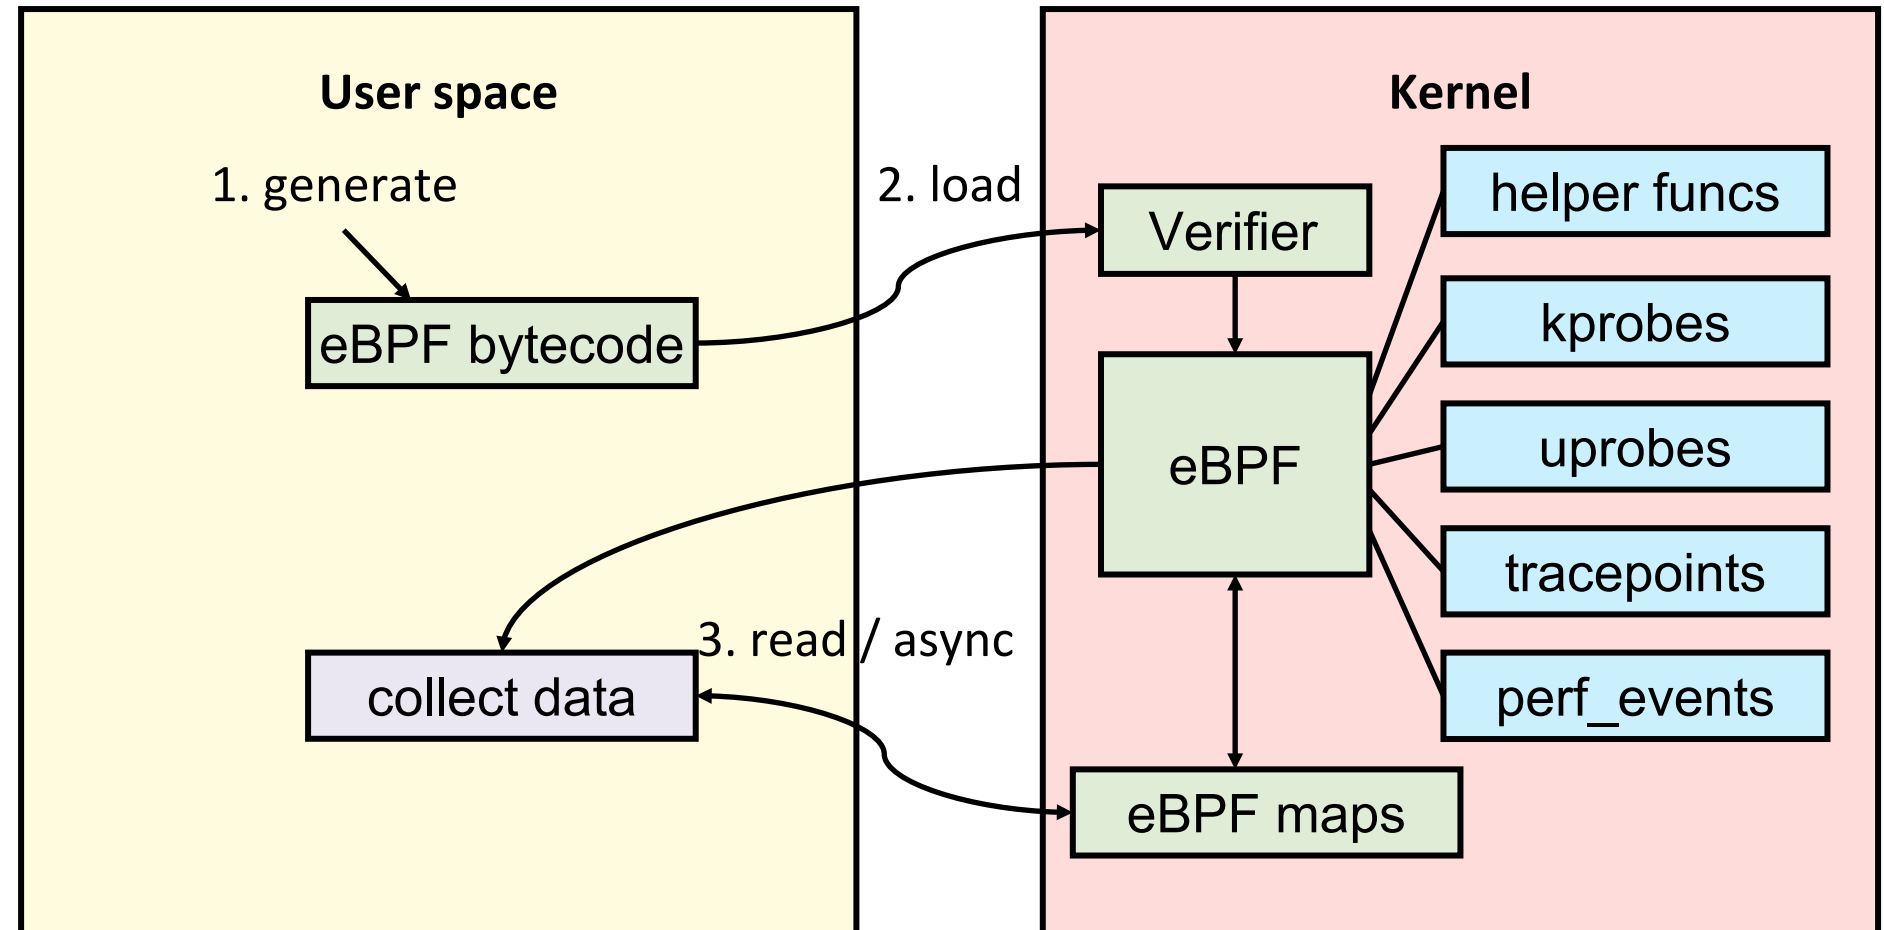


Agenda


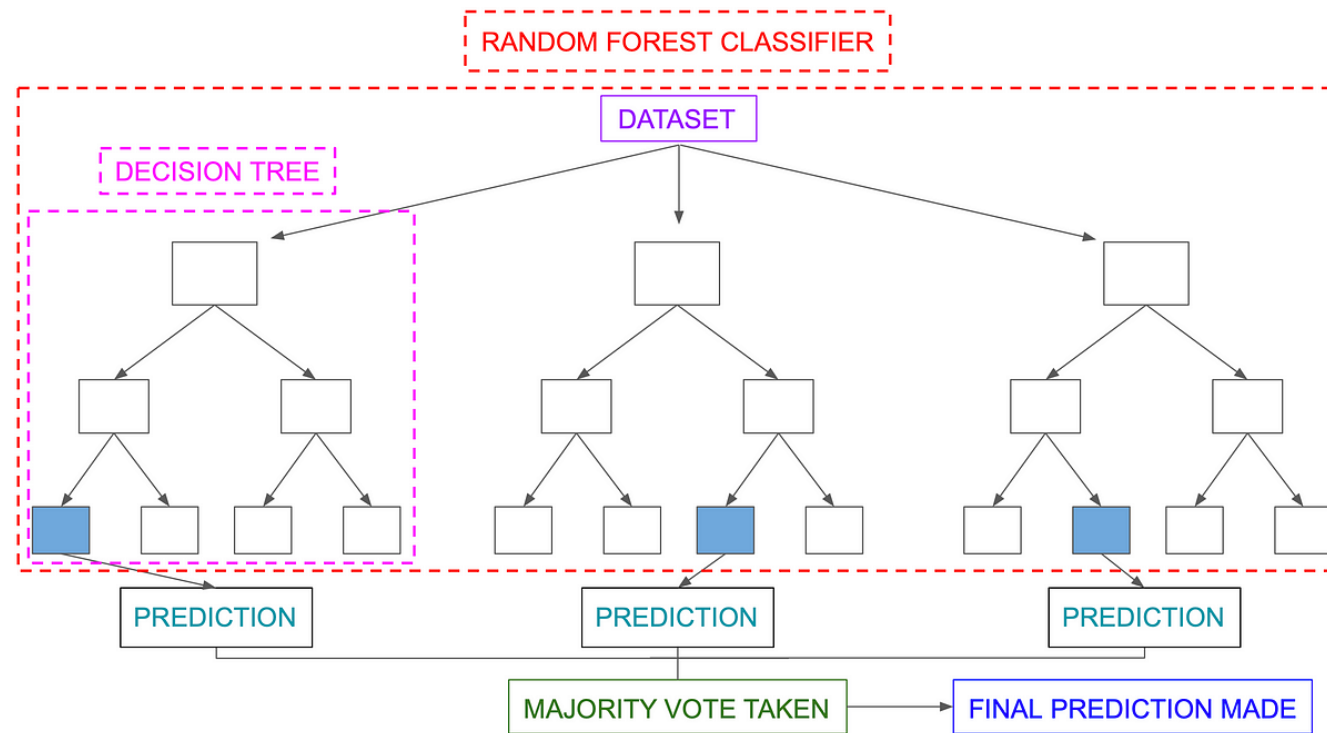
- Motivation
- Challenges
- **Overview & Technical Details**
 - Technical Backgrounds – eBPF & ML
 - O2Q Components
- Demo (CVE-2022-0995)
- Evaluation

Technical Background - eBPF

- Sandbox virtual machine in kernel.
- No need to modify kernel code or load module.
- Can hook any instruction.
- Own verifier.
- High performance using JIT.
- eBPF maps for data exchange.
- eBPF helper functions.

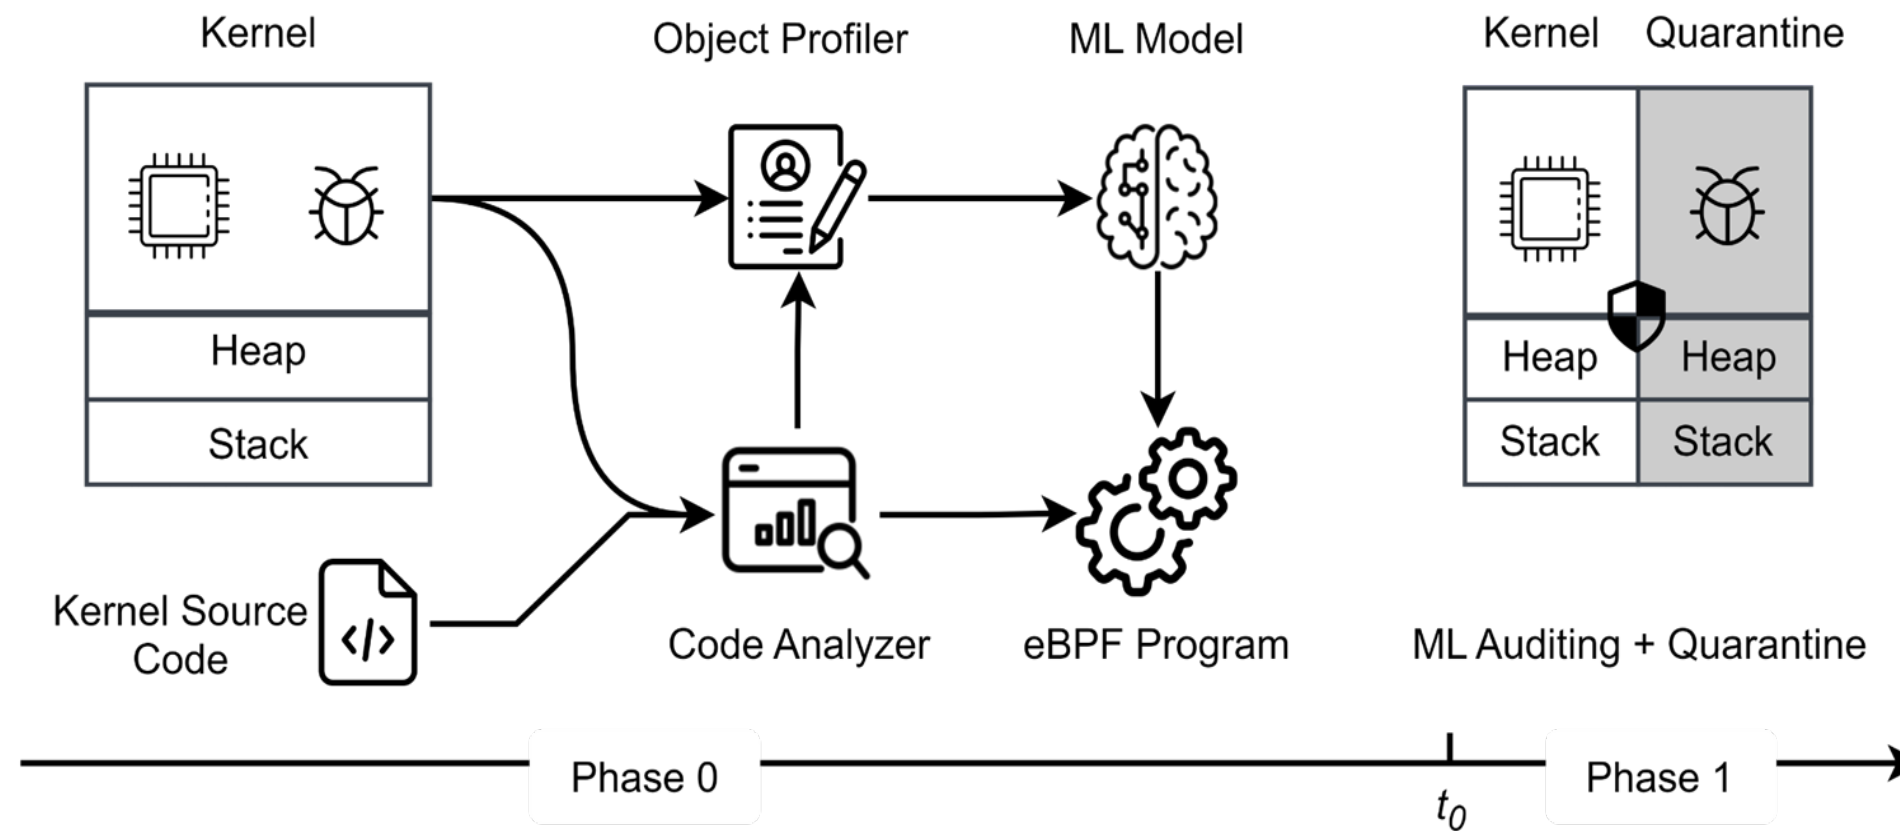


Technical Background - AI Models

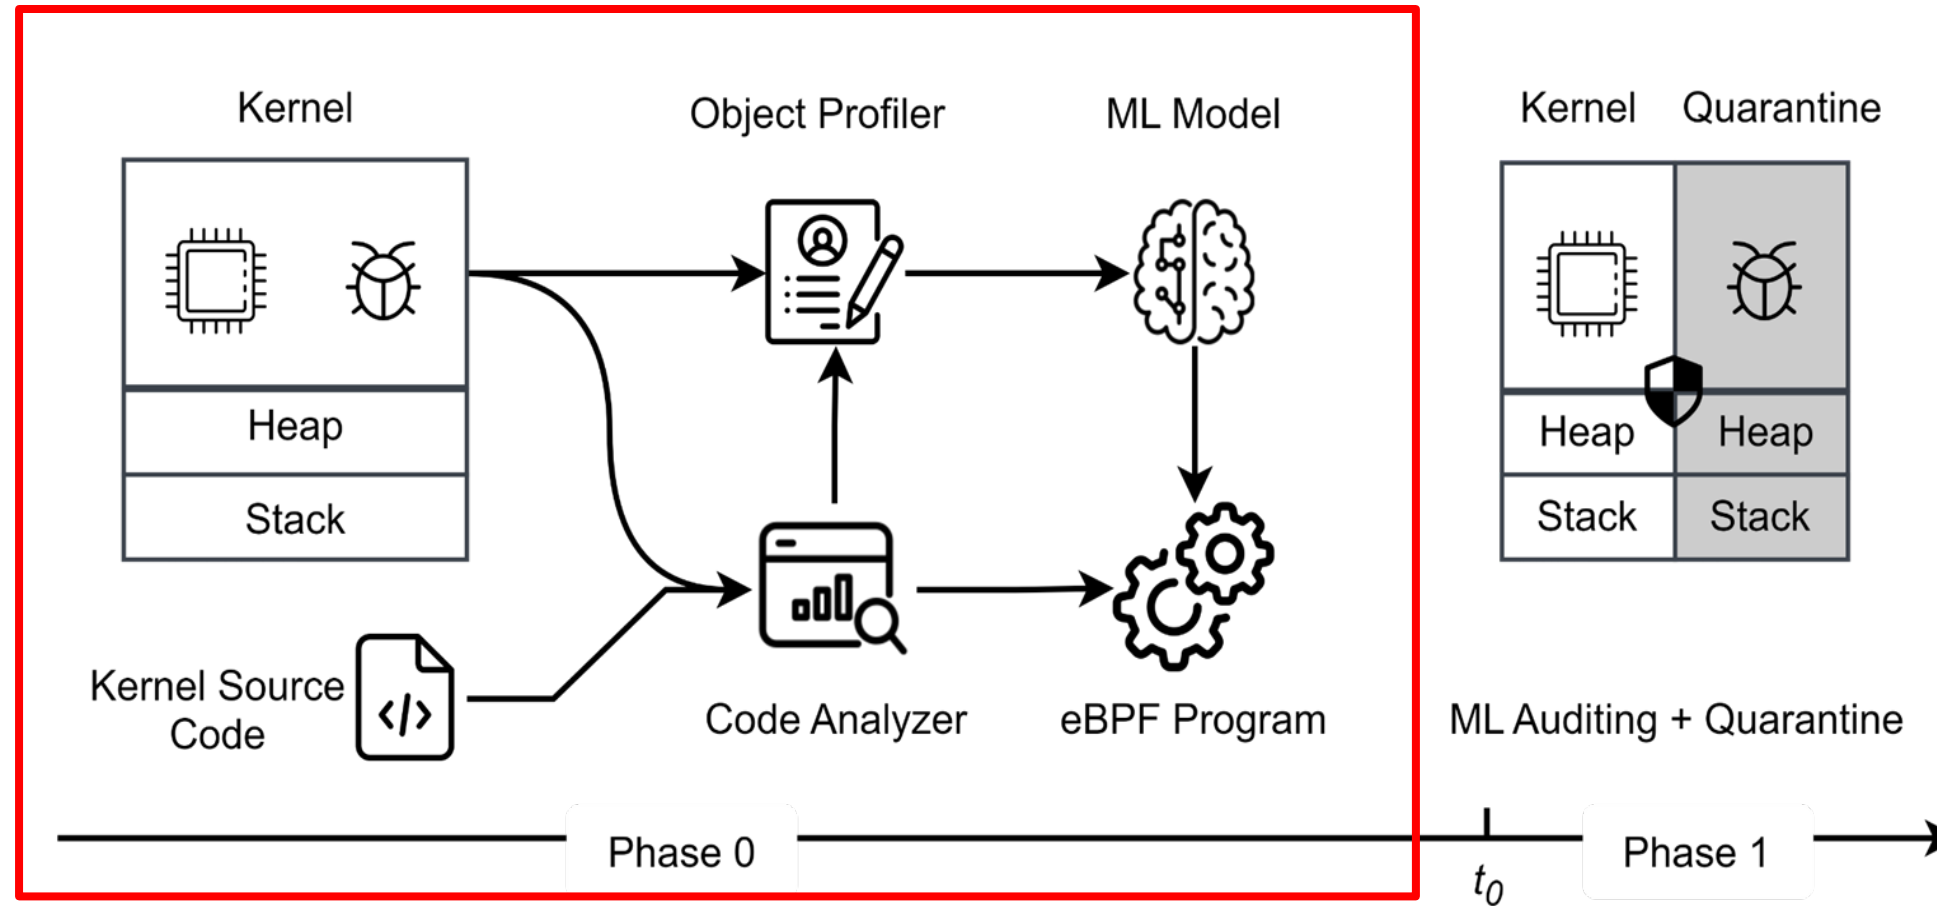


Source: <https://medium.com/analytics-vidhya/machine-learning-decision-trees-and-random-forest-classifiers-81422887a544>
<https://www.linkedin.com/pulse/introduction-neural-networks-how-machines-process-data-lakhani>

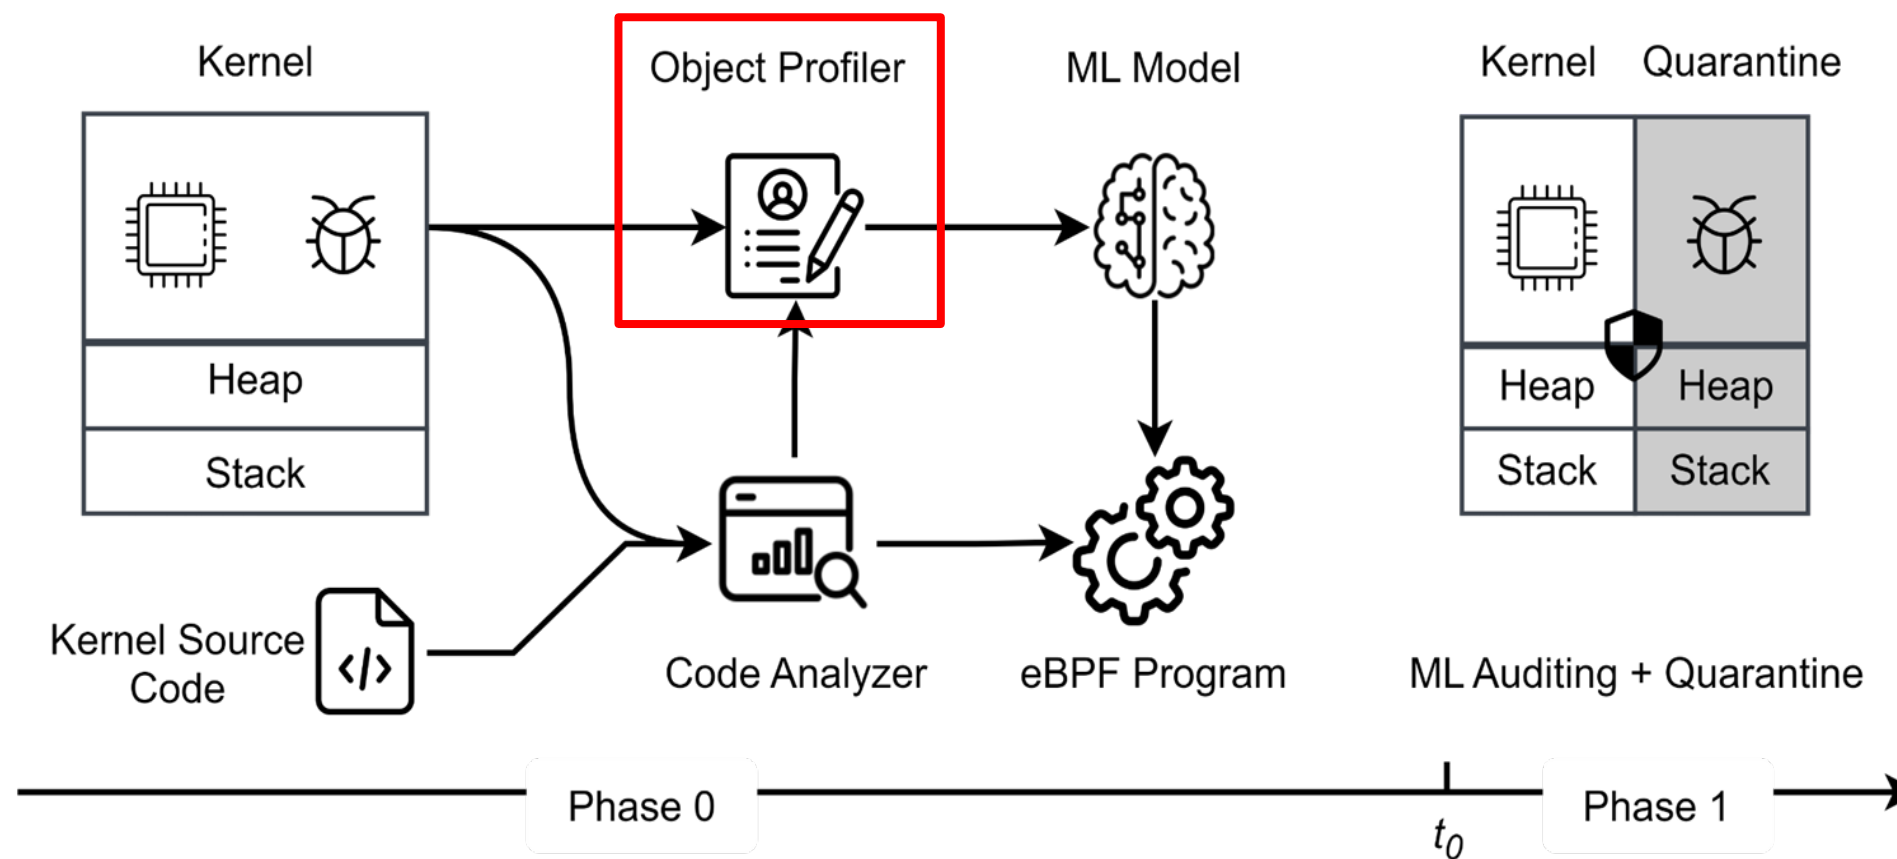
On-the-Fly Quarantine (O2Q) Overview



On-the-Fly Quarantine (O2Q) Overview

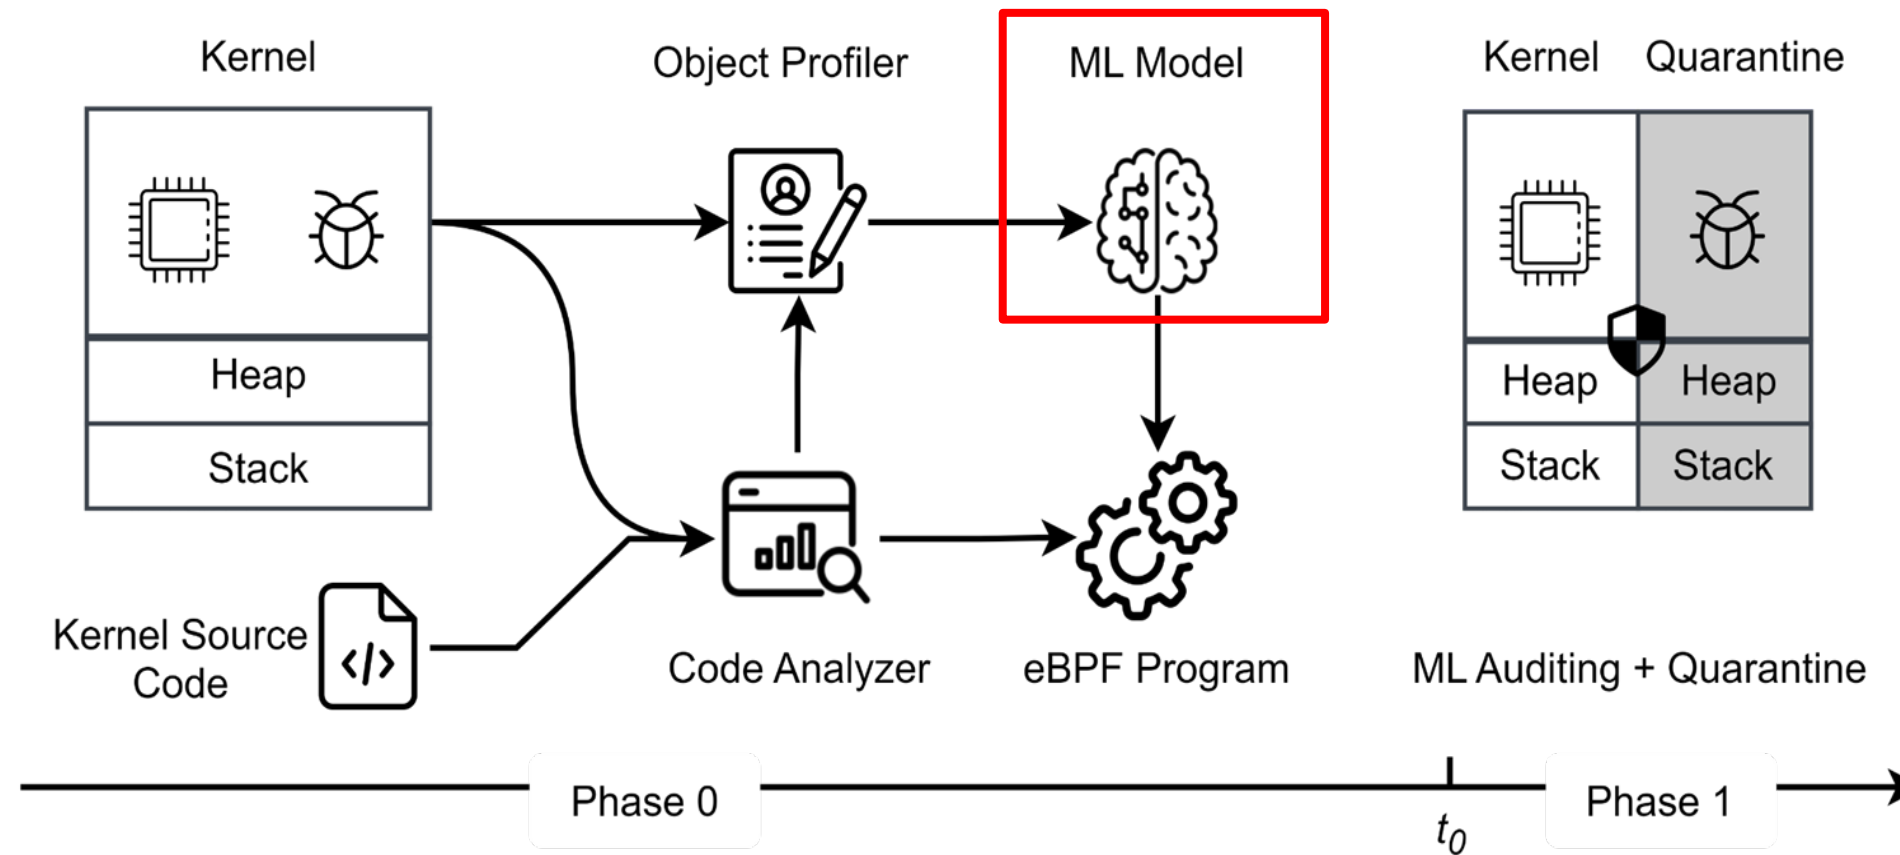


On-the-Fly Quarantine (O2Q) Overview



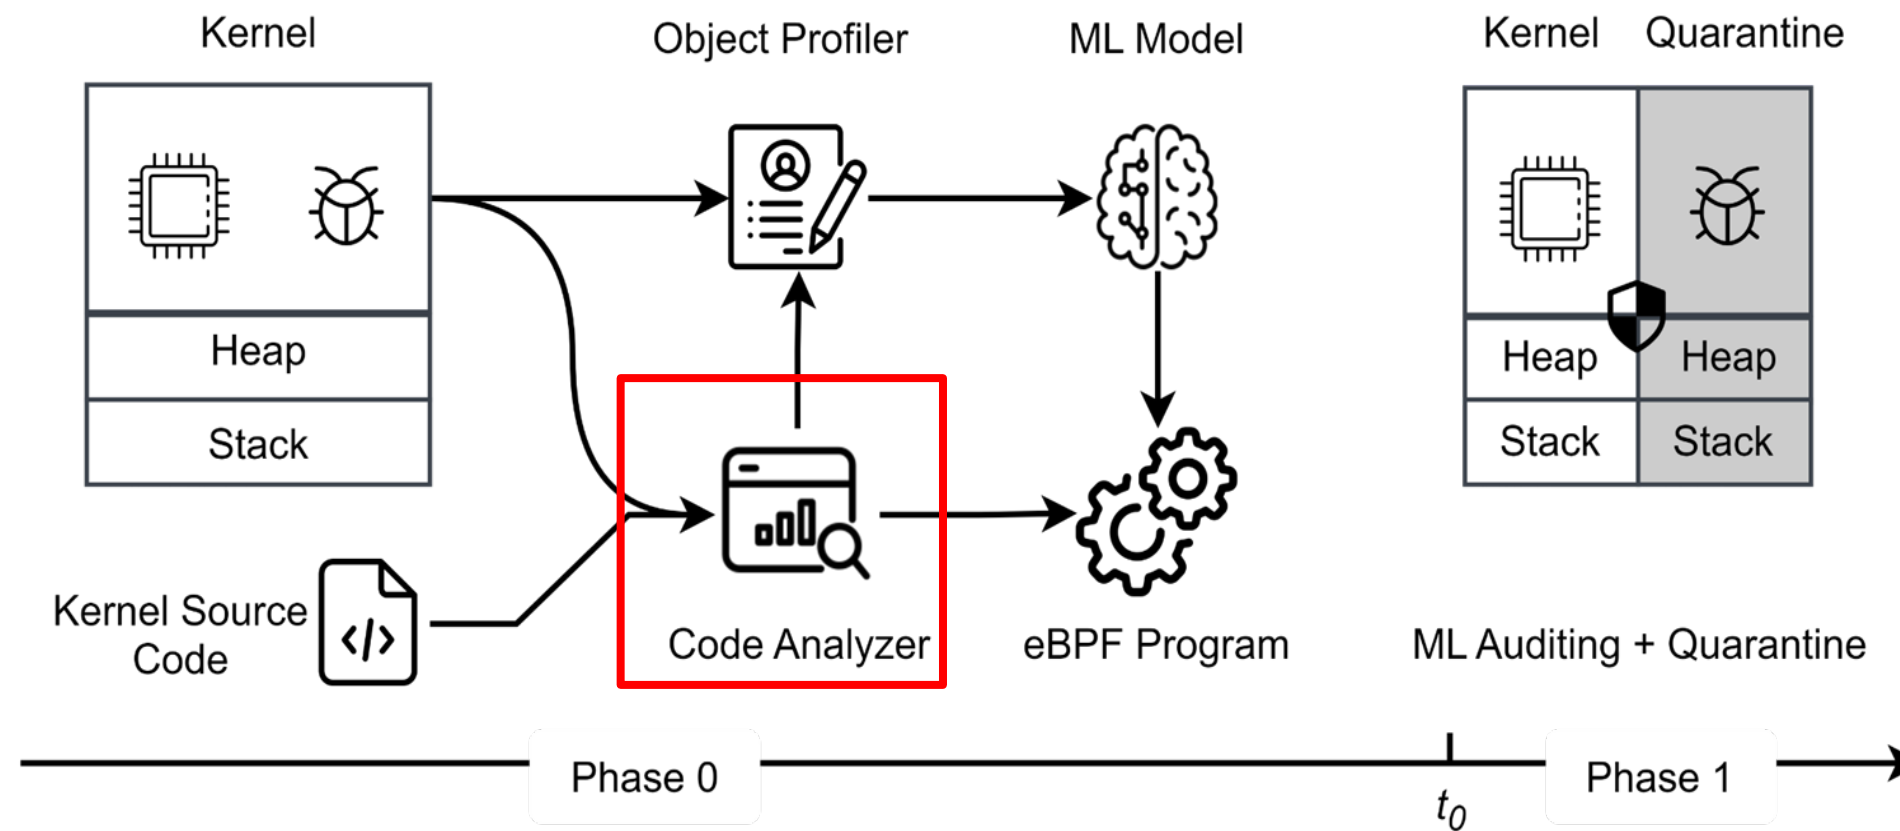
Collect data for ML model training: object's type and content

On-the-Fly Quarantine (O2Q) Overview



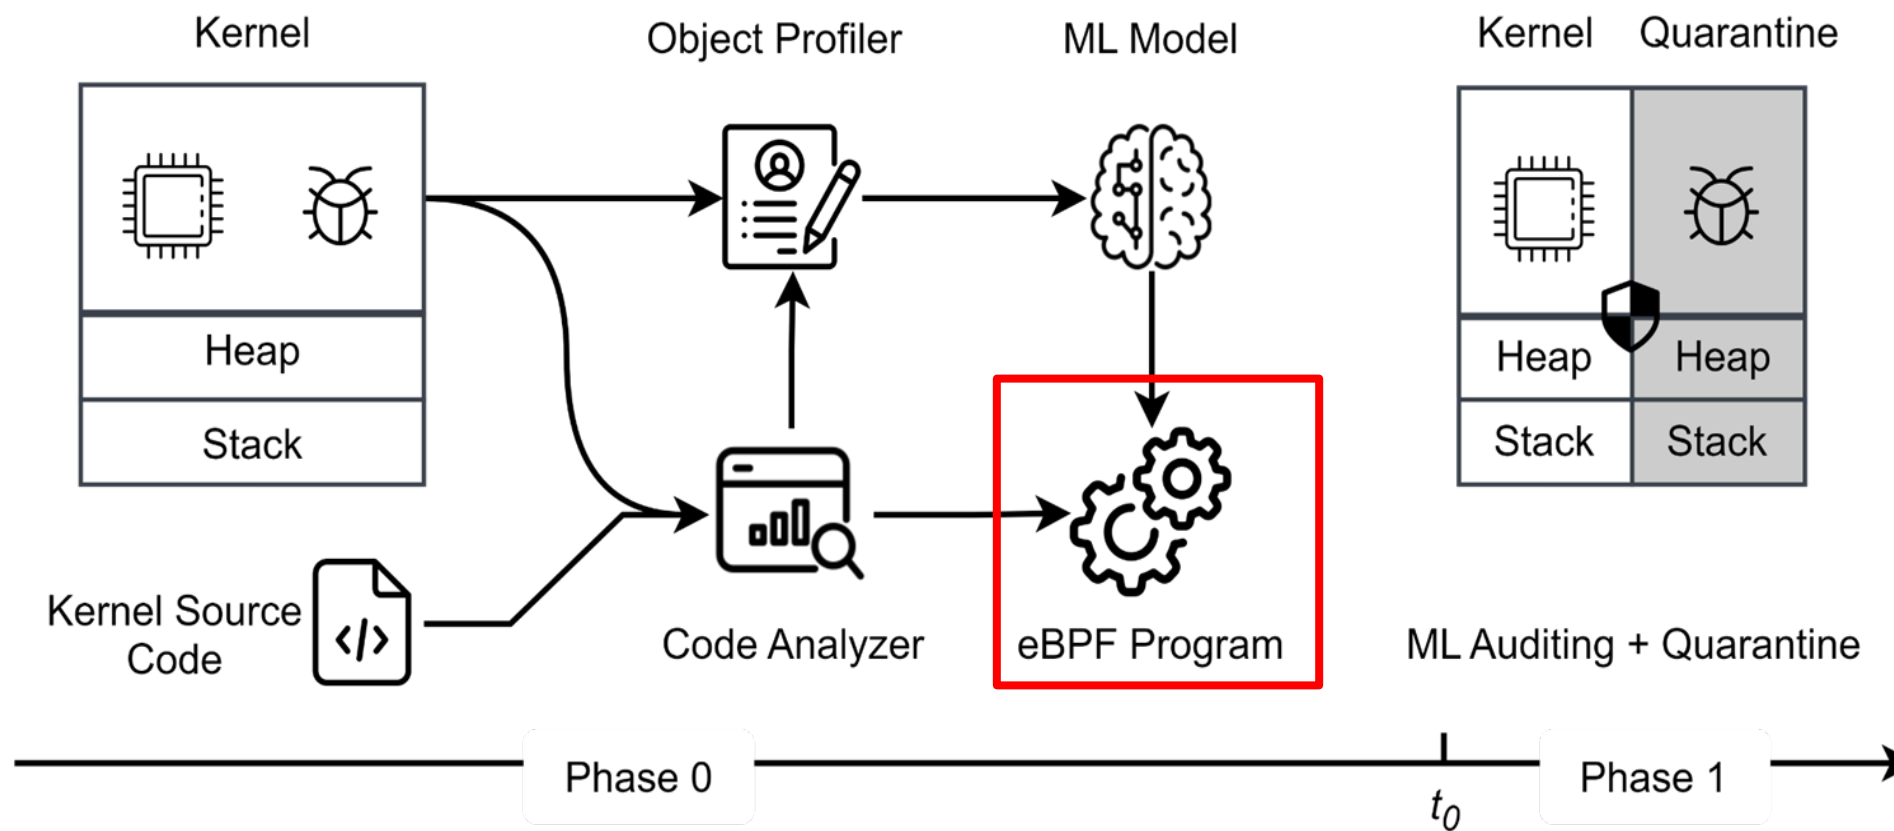
Train ML model inferring object's type based on its content

On-the-Fly Quarantine (O2Q) Overview



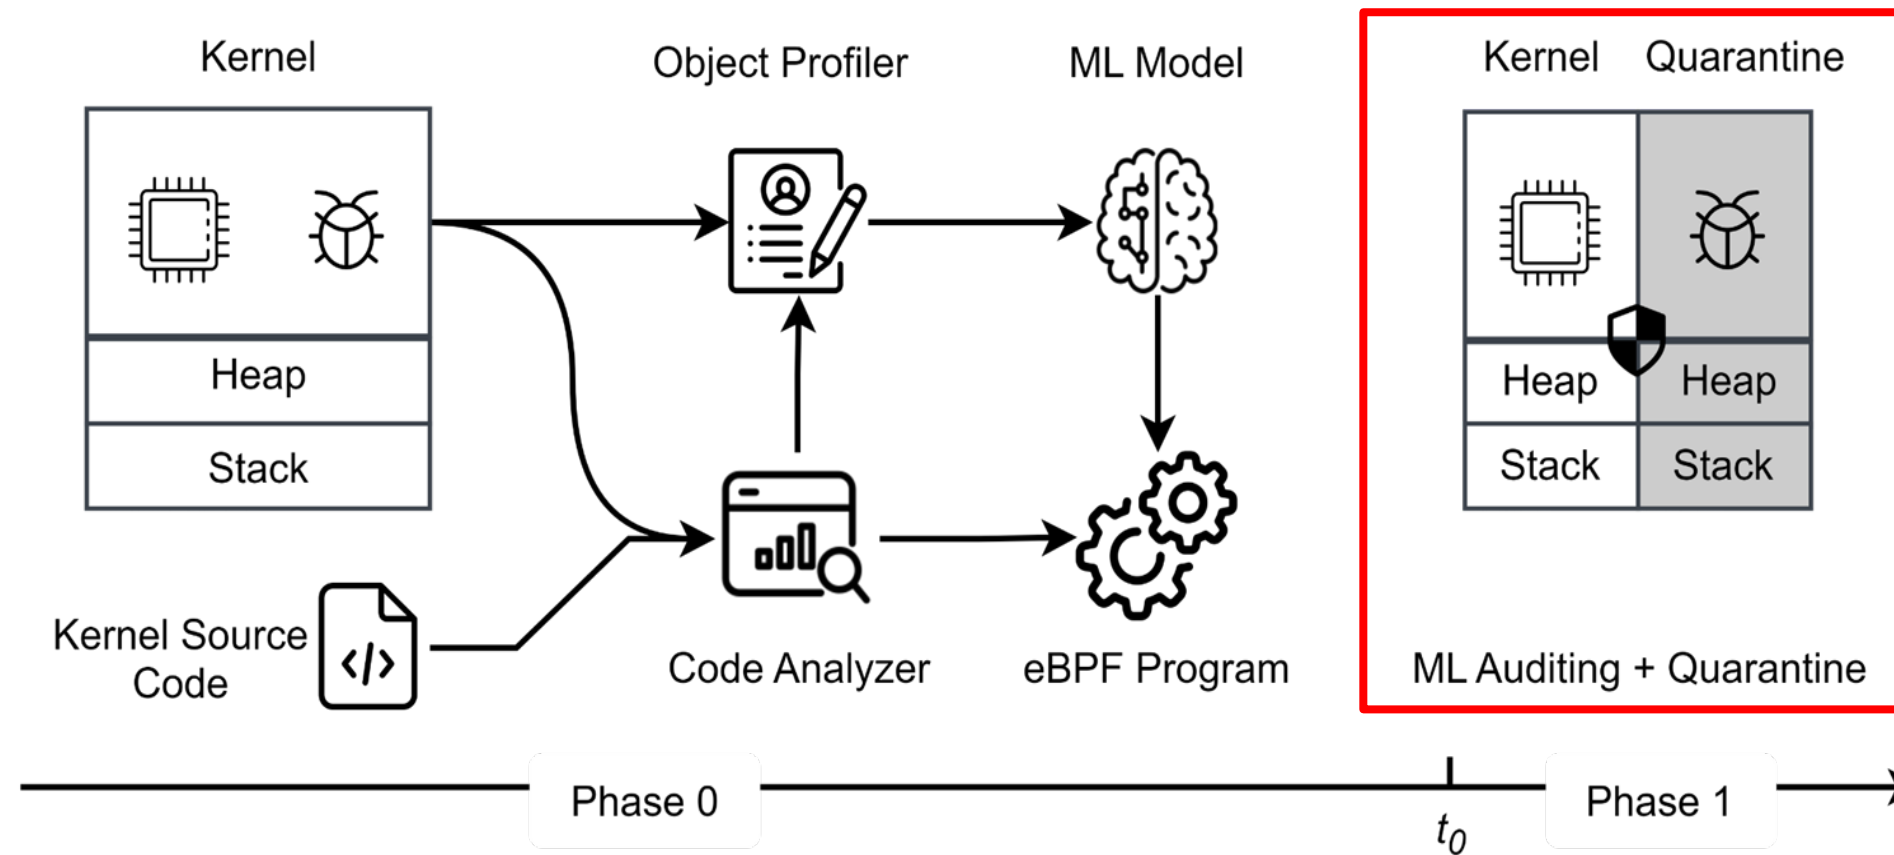
Identify instructions for instrumentation

On-the-Fly Quarantine (O2Q) Overview

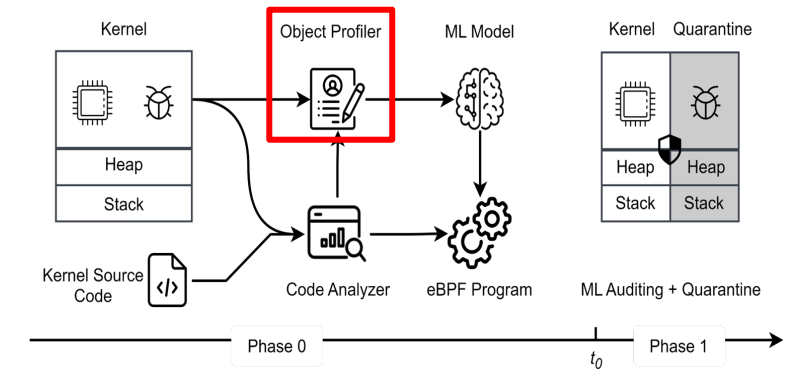


Implement quarantine, examine object's type at runtime

On-the-Fly Quarantine (O2Q) Overview

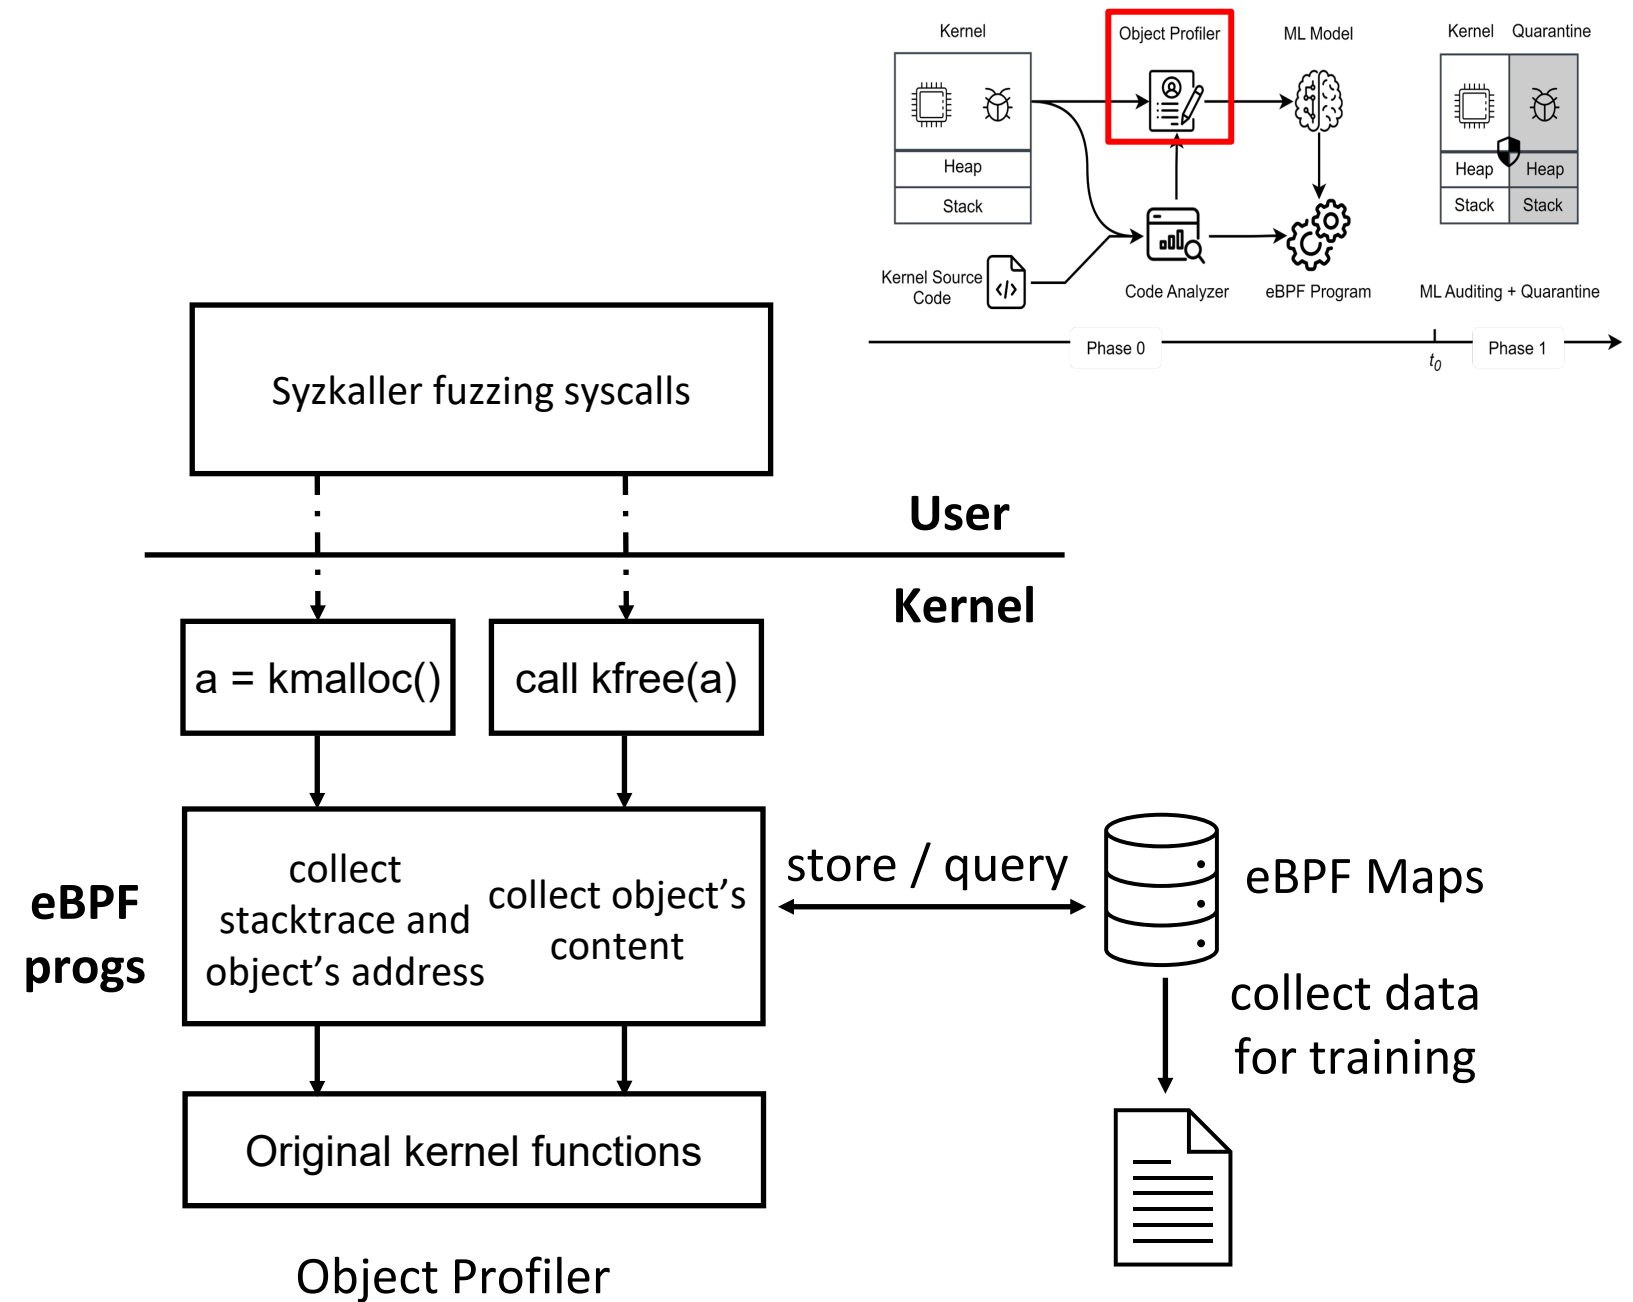


Object Profiler



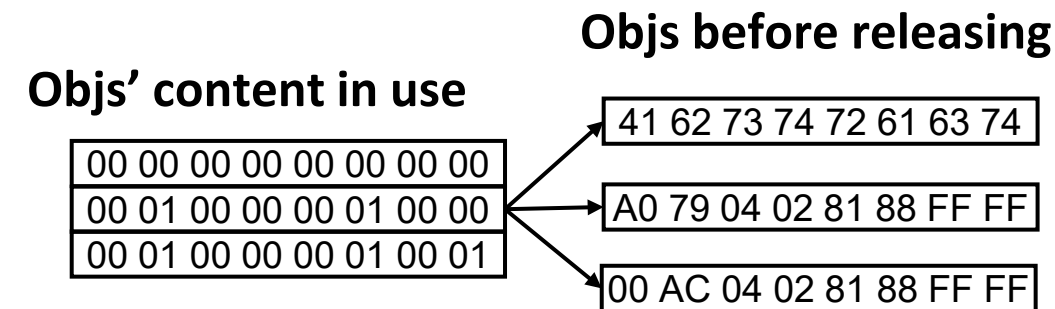
Object Profiler

- Use Syzkaller to enrich data source.
- Collect each object's content and type for training.

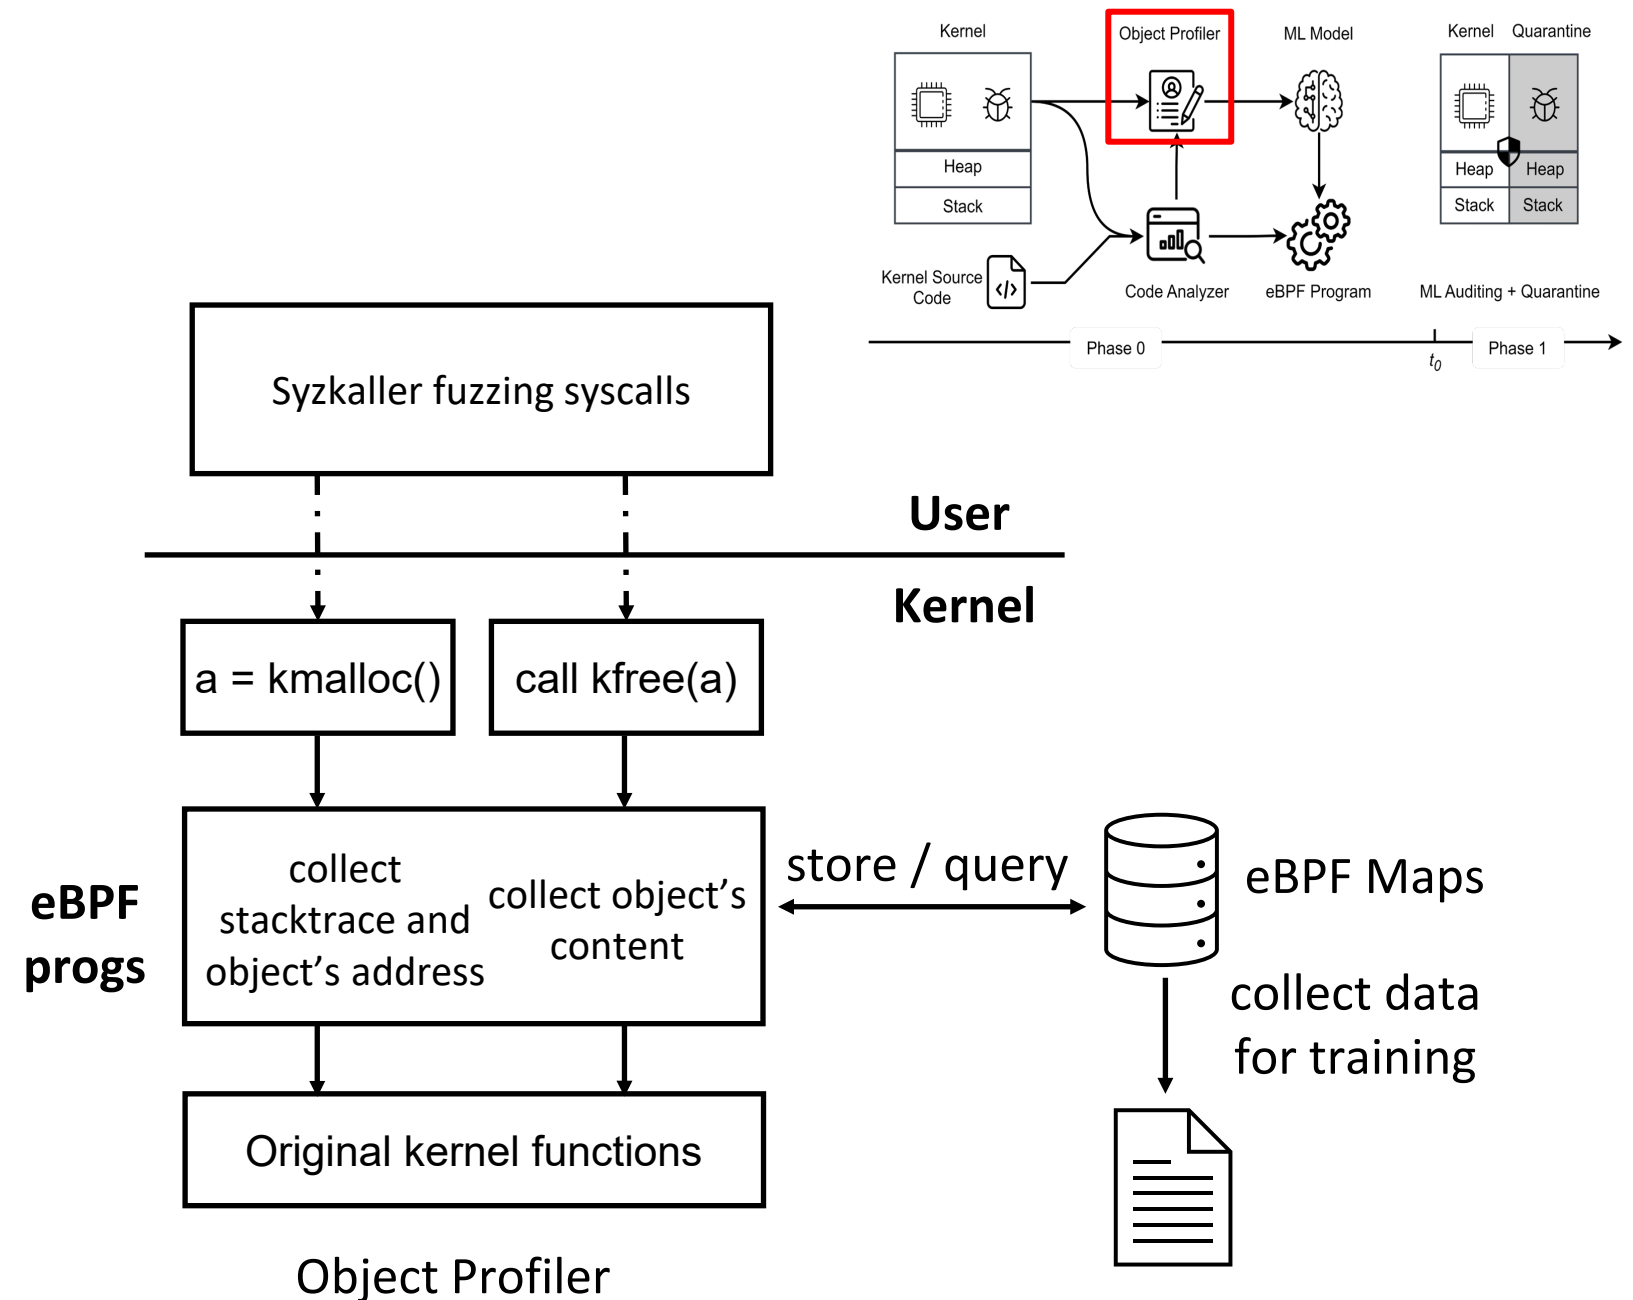


Object Profiler

- Use Syzkaller to enrich data source.
- Collect each object's content and type for training.
- Collect at object's release site: object possesses the most features that best reflect its characteristics.

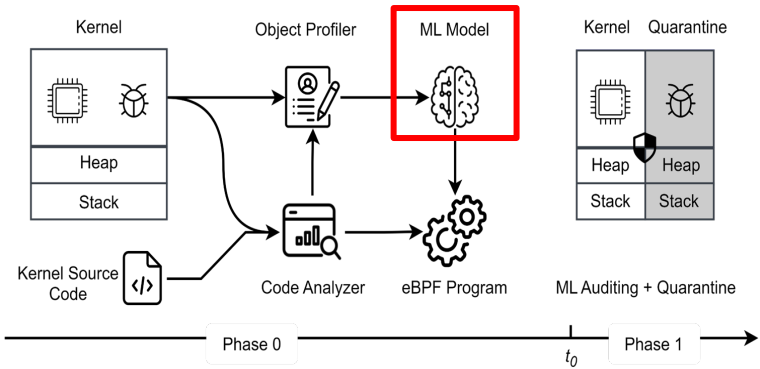


Uncharacterized vs. characterized



ML Model

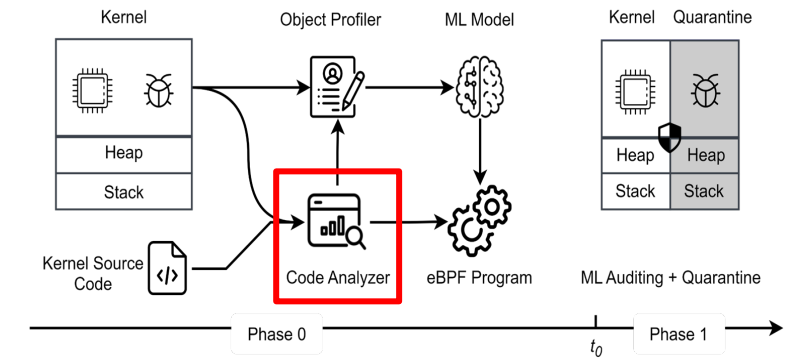
- Feature:
 - Object's data content as feature
- Label
 - Fine grained: object's type
 - Coarse grained: whether belongs to quarantine zone



	Tabular Data Processing	Interpretable	Defined Execution Time	Quantitative Accuracy	Convert to BPF Implementation
Decision Tree	✓	✓	✓	✓	✓
Random Forest	✓	✓	✓	✓	
Neural Network				✓	

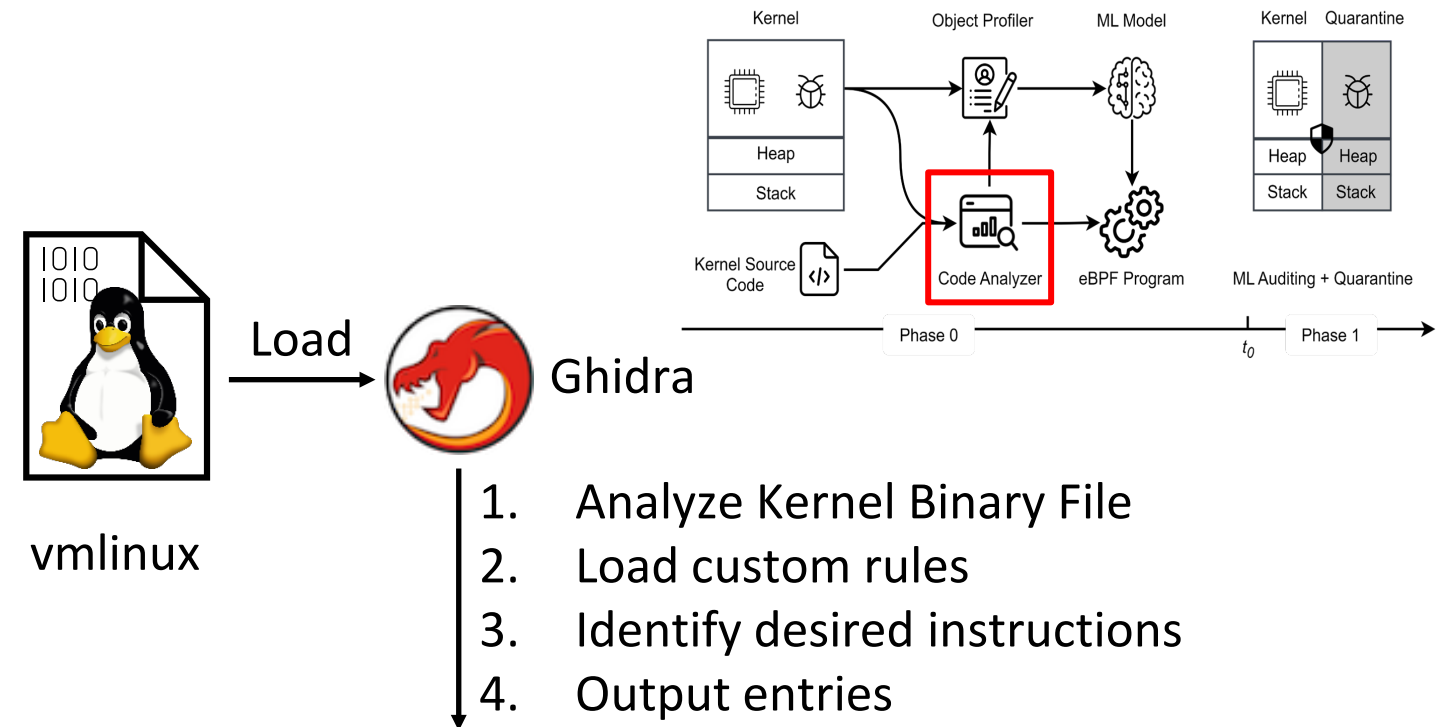
different ML model comparison

Code Analyzer



Code Analyzer

- Identify Linux Kernel's instructions
 - Indirect jump
 - Indirect call
 - Memory write
 - Subject switch



Indirect jump: `call *%rax`

Memory write: `mov $0x0, (%rsi, %rdx, 1)`

Determined address: `mov off(%rip), %rax`

Stack frame create: `sub offset, %rsp`

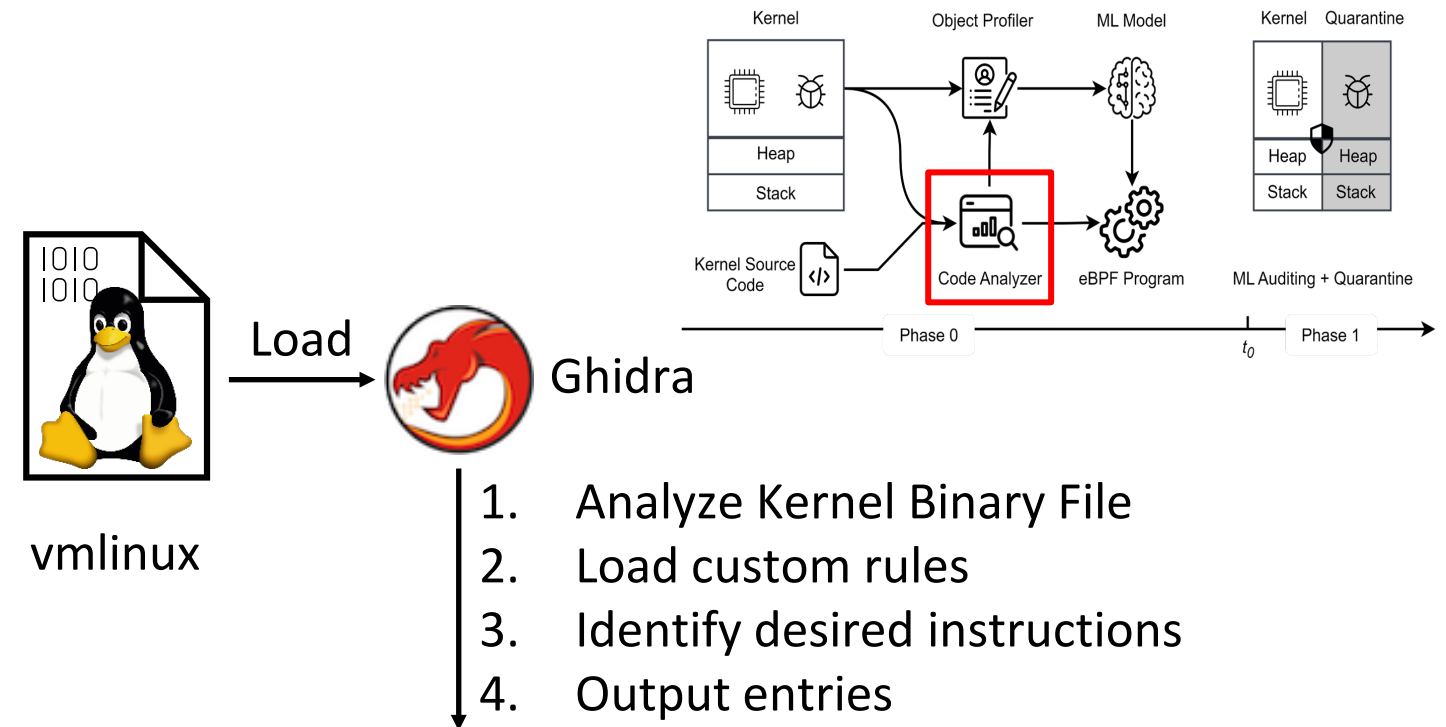
Stack access: `mov x, off (%rsp/rbp)`

Redundant check: `mov $0x0, off1(%rsi)`

Redundant check: `mov $0x0, off2(%rsi)`

Code Analyzer

- Identify Linux Kernel's instructions
 - Indirect jump
 - Indirect call
 - Memory write
 - Subject switch
- Efficiency Optimization:
 - Skip read
 - Skip determining address
 - Skip redundant check



Indirect jump: `call *%rax`

Memory write: `mov $0x0, (%rsi, %rdx, 1)`

Determined address: `mov off(%rip), %rax`

Stack frame create: `sub offset, %rsp`

Stack access: `mov x, off (%rsp/rbp)`


Redundant check: `mov $0x0, off1(%rsi)`

Redundant check: `mov $0x0, off2(%rsi)`

Code Analyzer

- Identify Linux Kernel's instructions
 - Indirect jump
 - Indirect call
 - Memory write
 - Subject switch
- Efficiency Optimization:
 - Skip read
 - Skip determining address
 - Skip redundant check

Reduced 24.07% instrument entries.



Indirect jump: `call *%rax`

Memory write: `mov $0x0, (%rsi, %rdx, 1)`

Determined address: `mov off(%rip), %rax`

Stack frame create: `sub offset, %rsp`

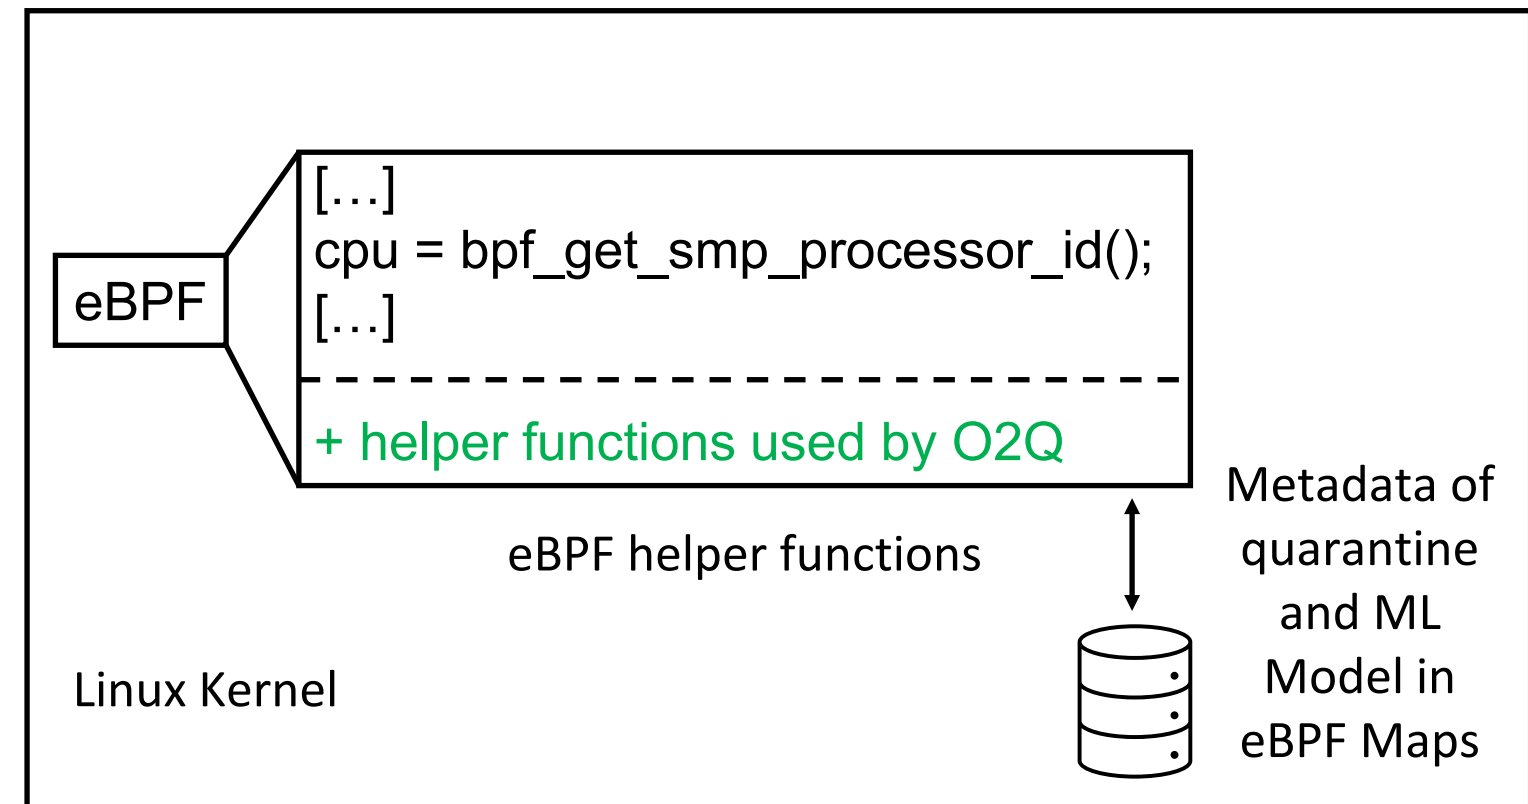
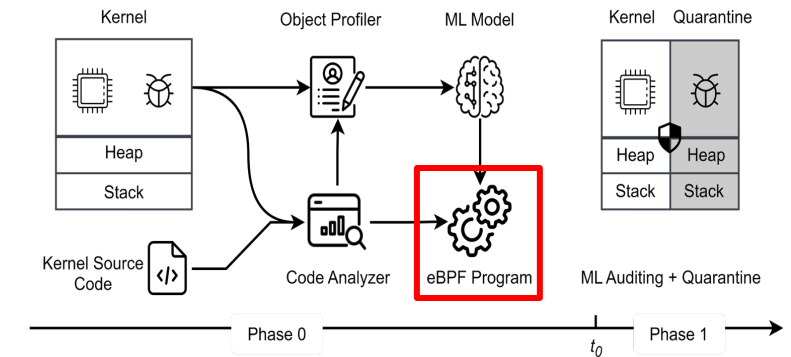
Stack access: `mov x, off (%rsp/rbp)`

Redundant check: `mov $0x0, off1(%rsi)`

Redundant check: `mov $0x0, off2(%rsi)`

eBPF Program

- Add extra eBPF helper functions for O2Q's functionality:
 - **bpf_set_regs()**: set register values, for switching stacks.
 - **bpf_create_slab_cache()**: creates private slab caches for the need of quarantine zone.
 - **bpf_cache_alloc() / bpf_cache_free()**: allocates from and frees to private caches.
- For better interaction with quarantine zone data:
 - **bpf_get_slab_*() / bpf_get_vm_struct()**: get the description of the slab and vmalloc directly, without traversing slab pages or vm_struct rb-trees.



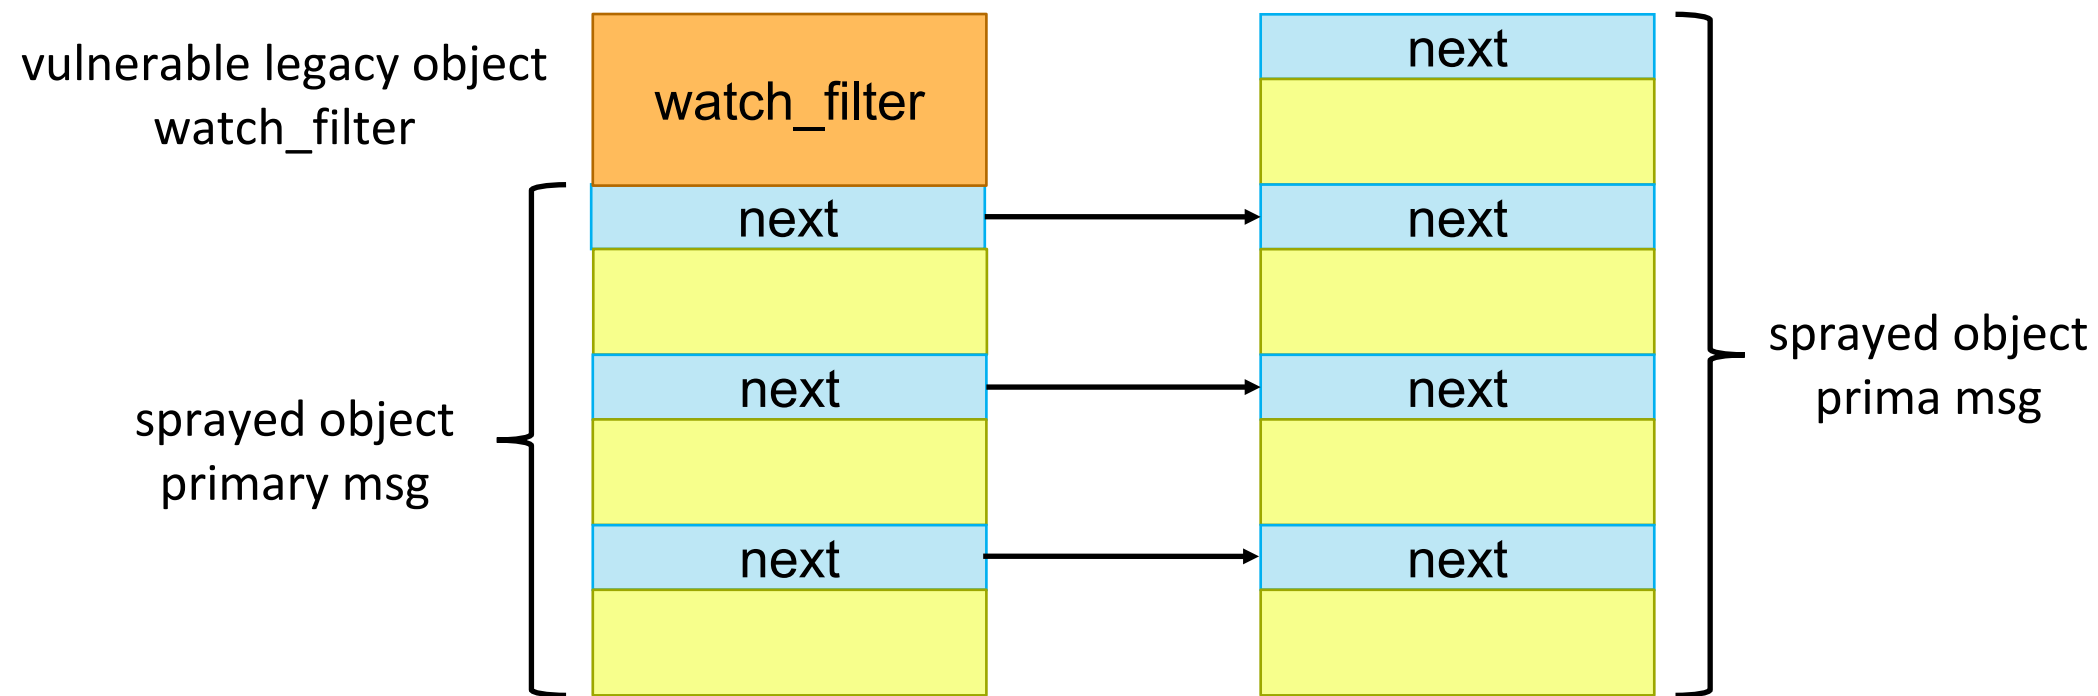
Agenda

- Motivation
- Challenges
- Overview & Technical Details
- **Demo (CVE-2022-0995)**
- Evaluation

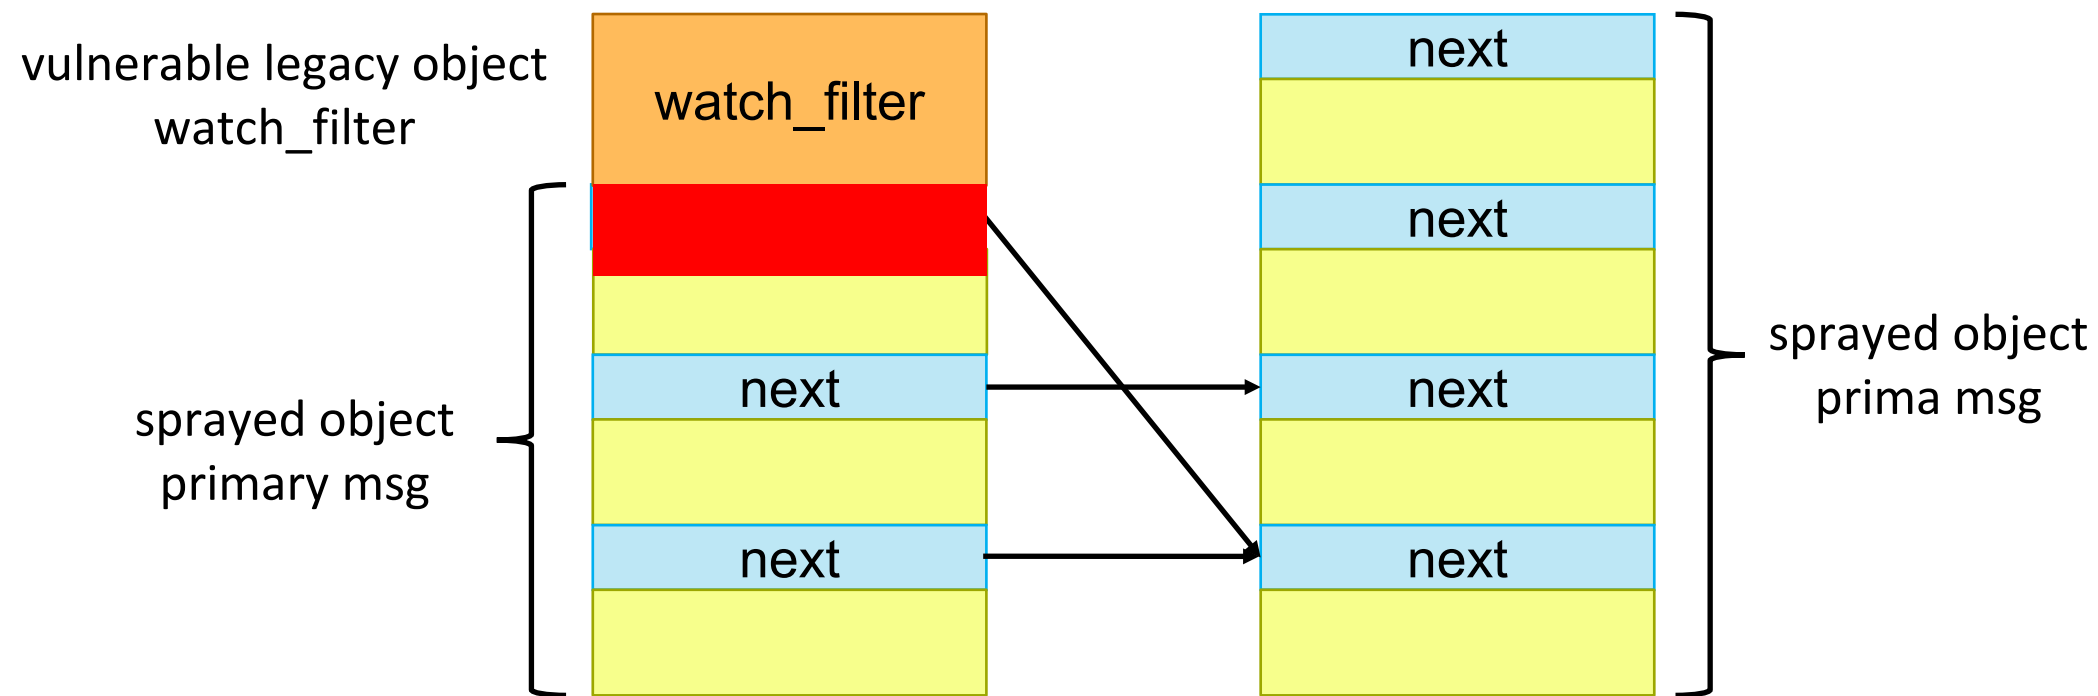


A Working Example: CVE-2022-0995

A Working Example: CVE-2022-0995



A Working Example: CVE-2022-0995



Two primary msg reference
this secondary msg.
Results in UAF

O2Q Workflow on CVE-2022-0995

O2Q Workflow on CVE-2022-0995

- The kernel is executing vulnerable component in quarantine zone.


Access



Untracked
memory


O2Q Workflow on CVE-2022-0995

- The kernel is executing vulnerable component in quarantine zone.
- The executing instruction should access watch_filter by Code Analyzer and Object Profiler.
- The eBPF program instrumented to the executing instructions encompasses the trained ML model.



O2Q Workflow on CVE-2022-0995

- The kernel is executing vulnerable component in quarantine zone.
- The executing instruction should access `watch_filter` by Code Analyzer and Object Profiler.
- The eBPF program instrumented to the executing instructions encompasses the trained ML model.
- The ML model infers the accessed object is `msg_msg`, indicating error.



```
test@syzkaller:~$
```

Original UAF exploits,
Two references to 8179

+



```
root@syzkaller:~/bpf# ./o2q_CVE-2022-0995
```

```
test@syzkaller:~$ uname -a
```

Deploy quarantine

Run exploit again, and
simulate legacy objects




Agenda


- Motivation
- Challenges
- Overview & Technical Details
- Demo (CVE-2022-0995)
- **Evaluation**

Evaluation

Overall system overhead



Evaluation



[3] HAKC: Preventing Kernel Hacks with HAKCs, NDSS'22

Evaluation

	Per Type		Per Component	
	Accuracy	Macro F1	Accuracy	Macro F1
IPV6				
Decision Tree	96.88 \pm 0.65	75.56 \pm 1.84	99.99 \pm 0.02	99.98 \pm 0.03
Random Forest	96.91 \pm 0.63	78.81 \pm 0.73	100 \pm 0.01	99.99 \pm 0.01
Neural Network	89.63 \pm 1.29	38.76 \pm 2.70	99.99 \pm 0.01	99.99 \pm 0.01
Sched				
Decision Tree	80.48 \pm 0.76	71.04 \pm 1.77	99.93 \pm 0.14	97.74 \pm 4.22
Random Forest	80.61 \pm 0.69	76.28 \pm 0.49	100 \pm 0	99.99 \pm 0.01
Neural Network	65.98 \pm 6.91	39.18 \pm 1.48	99.66 \pm 0.03	89.47 \pm 1.20
Netfilter				
Decision Tree	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Random Forest	89.54 \pm 0.15	81.87 \pm 1.86	99.96 \pm 0.05	99.77 \pm 0.29
Neural Network	72.9 \pm 2.23	37.98 \pm 2.83	97.16 \pm 0.17	74 \pm 2.56

performance of ML auditing

Evaluation

	Per Type		Per Component	
	Accuracy	Macro F1	Accuracy	Macro F1
IPV6				
Decision Tree	96.88 \pm 0.65	75.56 \pm 1.84	99.99 \pm 0.02	99.98 \pm 0.03
Random Forest	96.91 \pm 0.63	78.81 \pm 0.73	100 \pm 0.01	99.99 \pm 0.01
Neural Network	89.63 \pm 1.29	38.76 \pm 2.70	99.99 \pm 0.01	99.99 \pm 0.01
Sched				
Decision Tree	80.48 \pm 0.76	71.04 \pm 1.77	99.93 \pm 0.14	97.74 \pm 4.22
Random Forest	80.61 \pm 0.69	76.28 \pm 0.49	100 \pm 0	99.99 \pm 0.01
Neural Network	65.98 \pm 6.91	39.18 \pm 1.48	99.66 \pm 0.03	89.47 \pm 1.20
Netfilter				
Decision Tree	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Random Forest	89.54 \pm 0.15	81.87 \pm 1.86	99.96 \pm 0.05	99.77 \pm 0.29
Neural Network	72.9 \pm 2.23	37.98 \pm 2.83	97.16 \pm 0.17	74 \pm 2.56

performance of ML auditing

Evaluation

	Per Type		Per Component	
	Accuracy	Macro F1	Accuracy	Macro F1
IPV6				
Decision Tree	96.88 \pm 0.65	75.56 \pm 1.84	99.99 \pm 0.02	99.98 \pm 0.03
Random Forest	96.91 \pm 0.63	78.81 \pm 0.73	100 \pm 0.01	99.99 \pm 0.01
Neural Network	89.63 \pm 1.29	38.76 \pm 2.70	99.99 \pm 0.01	99.99 \pm 0.01
Sched				
Decision Tree	80.48 \pm 0.76	71.04 \pm 1.77	99.93 \pm 0.14	97.74 \pm 4.22
Random Forest	80.61 \pm 0.69	76.28 \pm 0.49	100 \pm 0	99.99 \pm 0.01
Neural Network	65.98 \pm 6.91	39.18 \pm 1.48	99.66 \pm 0.03	89.47 \pm 1.20
Netfilter				
Decision Tree	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Random Forest	89.54 \pm 0.15	81.87 \pm 1.86	99.96 \pm 0.05	99.77 \pm 0.29
Neural Network	72.9 \pm 2.23	37.98 \pm 2.83	97.16 \pm 0.17	74 \pm 2.56

performance of ML auditing

Evaluation

	Per Type		Per Component	
	Accuracy	Macro F1	Accuracy	Macro F1
IPV6				
Decision Tree	96.88 \pm 0.65	75.56 \pm 1.84	99.99 \pm 0.02	99.98 \pm 0.03
Random Forest	96.91 \pm 0.63	78.81 \pm 0.73	100 \pm 0.01	99.99 \pm 0.01
Neural Network	89.63 \pm 1.29	38.76 \pm 2.70	99.99 \pm 0.01	99.99 \pm 0.01
Sched				
Decision Tree	80.48 \pm 0.76	71.04 \pm 1.77	99.93 \pm 0.14	97.74 \pm 4.22
Random Forest	80.61 \pm 0.69	76.28 \pm 0.49	100 \pm 0	99.99 \pm 0.01
Neural Network	65.98 \pm 6.91	39.18 \pm 1.48	99.66 \pm 0.03	89.47 \pm 1.20
Netfilter				
Decision Tree	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Random Forest	89.54 \pm 0.15	81.87 \pm 1.86	99.96 \pm 0.05	99.77 \pm 0.29
Neural Network	72.9 \pm 2.23	37.98 \pm 2.83	97.16 \pm 0.17	74 \pm 2.56

performance of ML auditing

Evaluation

	Per Type		Per Component	
	Accuracy	Macro F1	Accuracy	Macro F1
IPV6				
Decision Tree	96.88 \pm 0.65	75.56 \pm 1.84	99.99 \pm 0.02	99.98 \pm 0.03
Random Forest	96.91 \pm 0.63	78.81 \pm 0.73	100 \pm 0.01	99.99 \pm 0.01
Neural Network	89.63 \pm 1.29	38.76 \pm 2.70	99.99 \pm 0.01	99.99 \pm 0.01
Sched				
Decision Tree	80.48 \pm 0.76	71.04 \pm 1.77	99.93 \pm 0.14	97.74 \pm 4.22
Random Forest	80.61 \pm 0.69	76.28 \pm 0.49	100 \pm 0	99.99 \pm 0.01
Neural Network	65.98 \pm 6.91	39.18 \pm 1.48	99.66 \pm 0.03	89.47 \pm 1.20
Netfilter				
Decision Tree	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Random Forest				99.77 \pm 0.29

Neural Network is not good enough

Random Forest is too heavy to be embedded via eBPF

performance of ML auditing

Evaluation

	Accuracy	Macro F1	Accuracy	Macro F1
Feature Length				
32	88.40 \pm 0.42	73.97 \pm 3.83	98.75 \pm 0.41	91.91 \pm 2.32
64	89.15 \pm 0.33	77.24 \pm 4.21	99.91 \pm 0.07	99.47 \pm 0.45
128	89.18 \pm 0.29	77.44 \pm 4.33	99.85 \pm 0.1	99.46 \pm 0.64
256	89.26 \pm 0.29	77.34 \pm 5.06	99.92 \pm 0.08	99.51 \pm 0.49
1024	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Max Depth				
3	61.18 \pm 2.45	1.72 \pm 0.19	97.47 \pm 0.4	79.34 \pm 3.03
7	76.59 \pm 2.38	8.48 \pm 0.58	99.44 \pm 0.21	96.44 \pm 1.32
10	83.54 \pm 2.19	21.06 \pm 2.19	99.65 \pm 0.14	97.78 \pm 0.86
14	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46

Performance of tuning decision tree feature length and depth

Evaluation

	Accuracy	Macro F1	Accuracy	Macro F1
Feature Length				
32	88.40 \pm 0.42	73.97 \pm 3.83	98.75 \pm 0.41	91.91 \pm 2.32
64	89.15 \pm 0.33	77.24 \pm 4.21	99.91 \pm 0.07	99.47 \pm 0.45
128	89.18 \pm 0.29	77.44 \pm 4.33	99.85 \pm 0.1	99.46 \pm 0.64
256	89.26 \pm 0.29	77.34 \pm 5.06	99.92 \pm 0.08	99.51 \pm 0.49
1024	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Max Depth				
3	61.18 \pm 2.45	1.72 \pm 0.19	97.47 \pm 0.4	79.34 \pm 3.03
7	76.59 \pm 2.38	8.48 \pm 0.58	99.44 \pm 0.21	96.44 \pm 1.32
10	83.54 \pm 2.19	21.06 \pm 2.19	99.65 \pm 0.14	97.78 \pm 0.86
14	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46

Performance of tuning decision tree feature length and depth


Evaluation

	Accuracy	Macro F1	Accuracy	Macro F1
Feature Length				
32	88.40 \pm 0.42	73.97 \pm 3.83	98.75 \pm 0.41	91.91 \pm 2.32
64	89.15 \pm 0.33	77.24 \pm 4.21	99.91 \pm 0.07	99.47 \pm 0.45
128	89.18 \pm 0.29	77.44 \pm 4.33	99.85 \pm 0.1	99.46 \pm 0.64
256	89.26 \pm 0.29	77.34 \pm 5.06	99.92 \pm 0.08	99.51 \pm 0.49
1024	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46
Max Depth				
3	61.18 \pm 2.45	1.72 \pm 0.19	97.47 \pm 0.4	79.34 \pm 3.03
7	76.59 \pm 2.38	8.48 \pm 0.58	99.44 \pm 0.21	96.44 \pm 1.32
10	83.54 \pm 2.19	21.06 \pm 2.19	99.65 \pm 0.14	97.78 \pm 0.86
14	89.47 \pm 0.23	78.17 \pm 4.88	99.92 \pm 0.07	99.51 \pm 0.46

Performance of tuning decision tree feature length and depth

Takeaway

- Our work revealed the legacy object problem, which is critical to protect the kernel on-the-fly before patches are available.
- We demonstrated how embedding machine learning into the kernel can effectively solve the legacy object problem.
- Limitation: ML model accuracy is not 100%, only sufficing as a temporary remediation before patches are available.
- Future work:
 - Mature the prototype implementation and solution to corner cases in ML model. Expecting collaboration.
 - Reduce overhead using PKS like hardware feature.



Qinrun Dai (@2st___)

Looking for

2025 summer internship!